



# SFS: Smart OS Scheduling for Serverless Functions

Yuqi Fu<sup>1</sup>, Li Liu<sup>2</sup>, Haoliang Wang<sup>3</sup>, Yue Cheng<sup>1</sup>, Songqing Chen<sup>2</sup>

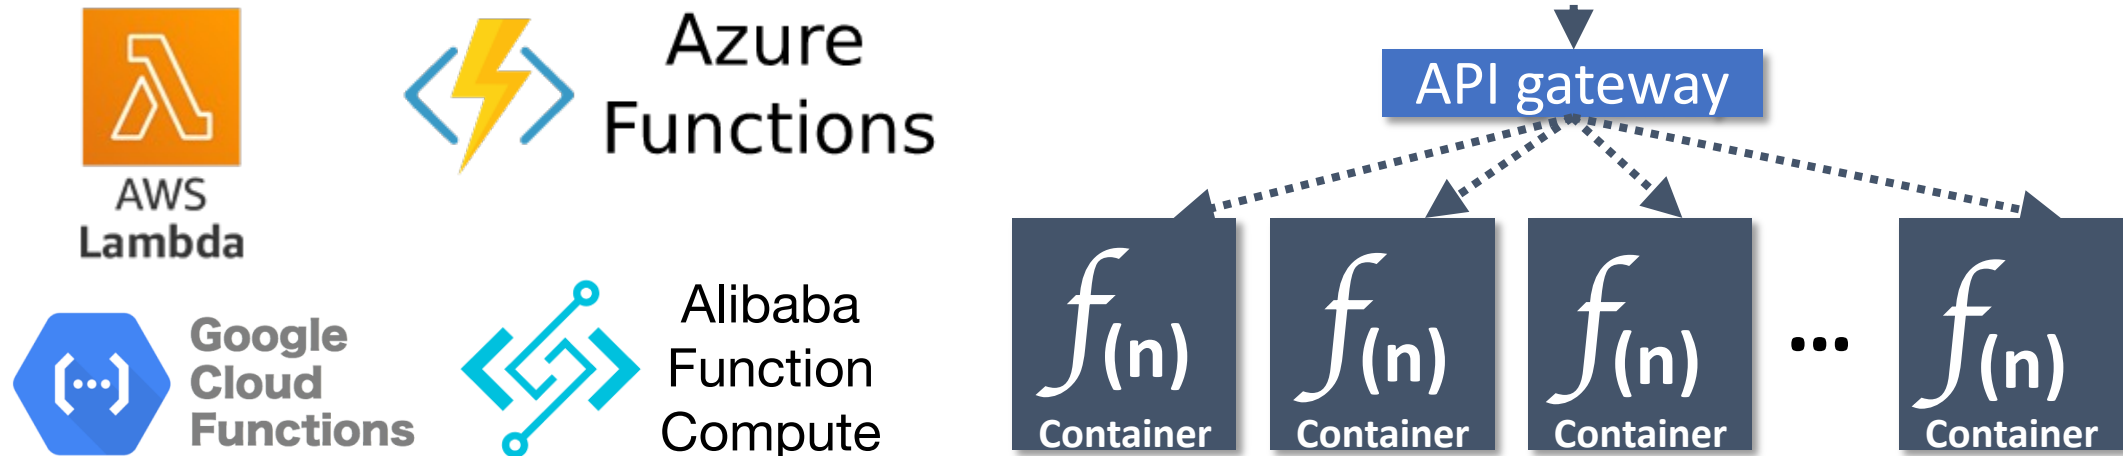
<sup>1</sup>University of Virginia, <sup>2</sup>George Mason University, <sup>3</sup>Adobe Research



# Serverless computing

A programming abstraction that enables users to upload programs, run them at **virtually** any scale, and pay **only for the resources used**

**Function-as-a-Service (FaaS):** Cloud functions as a basic deployment unit



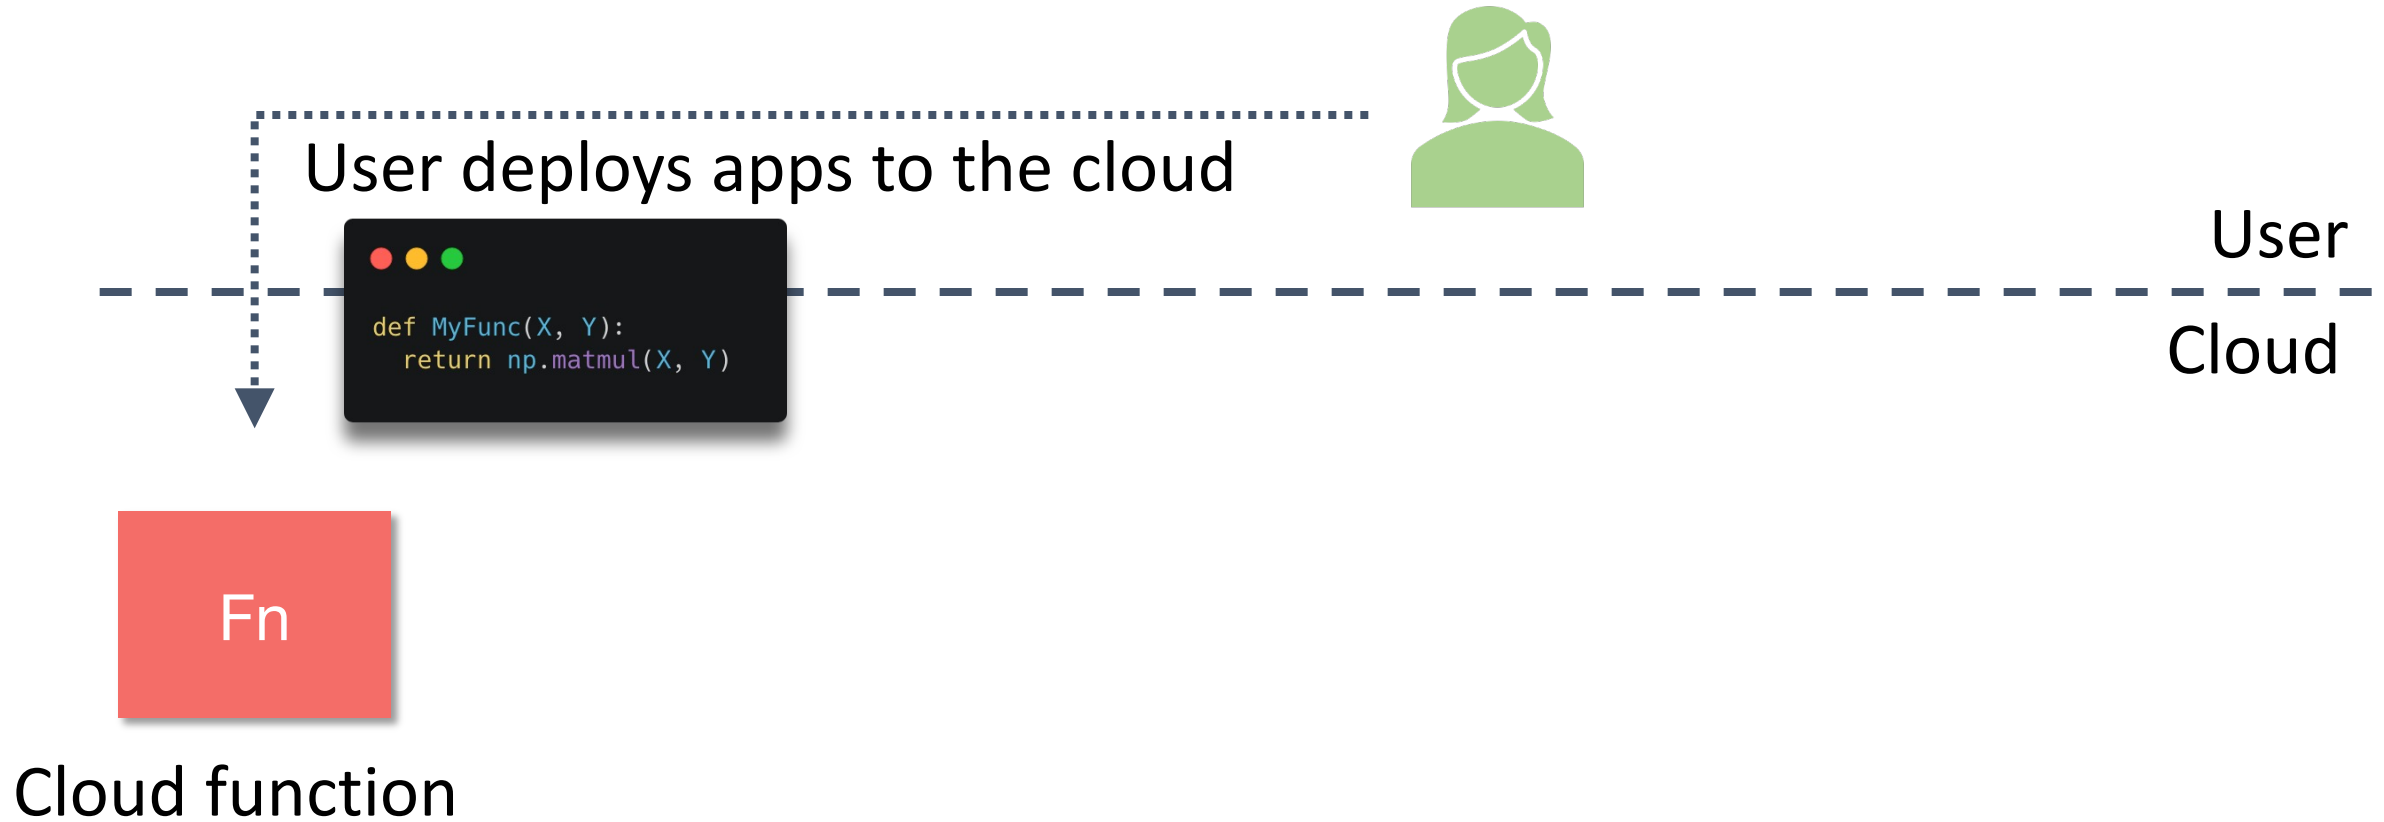
# Function-as-a-Service (FaaS)




User  
Cloud



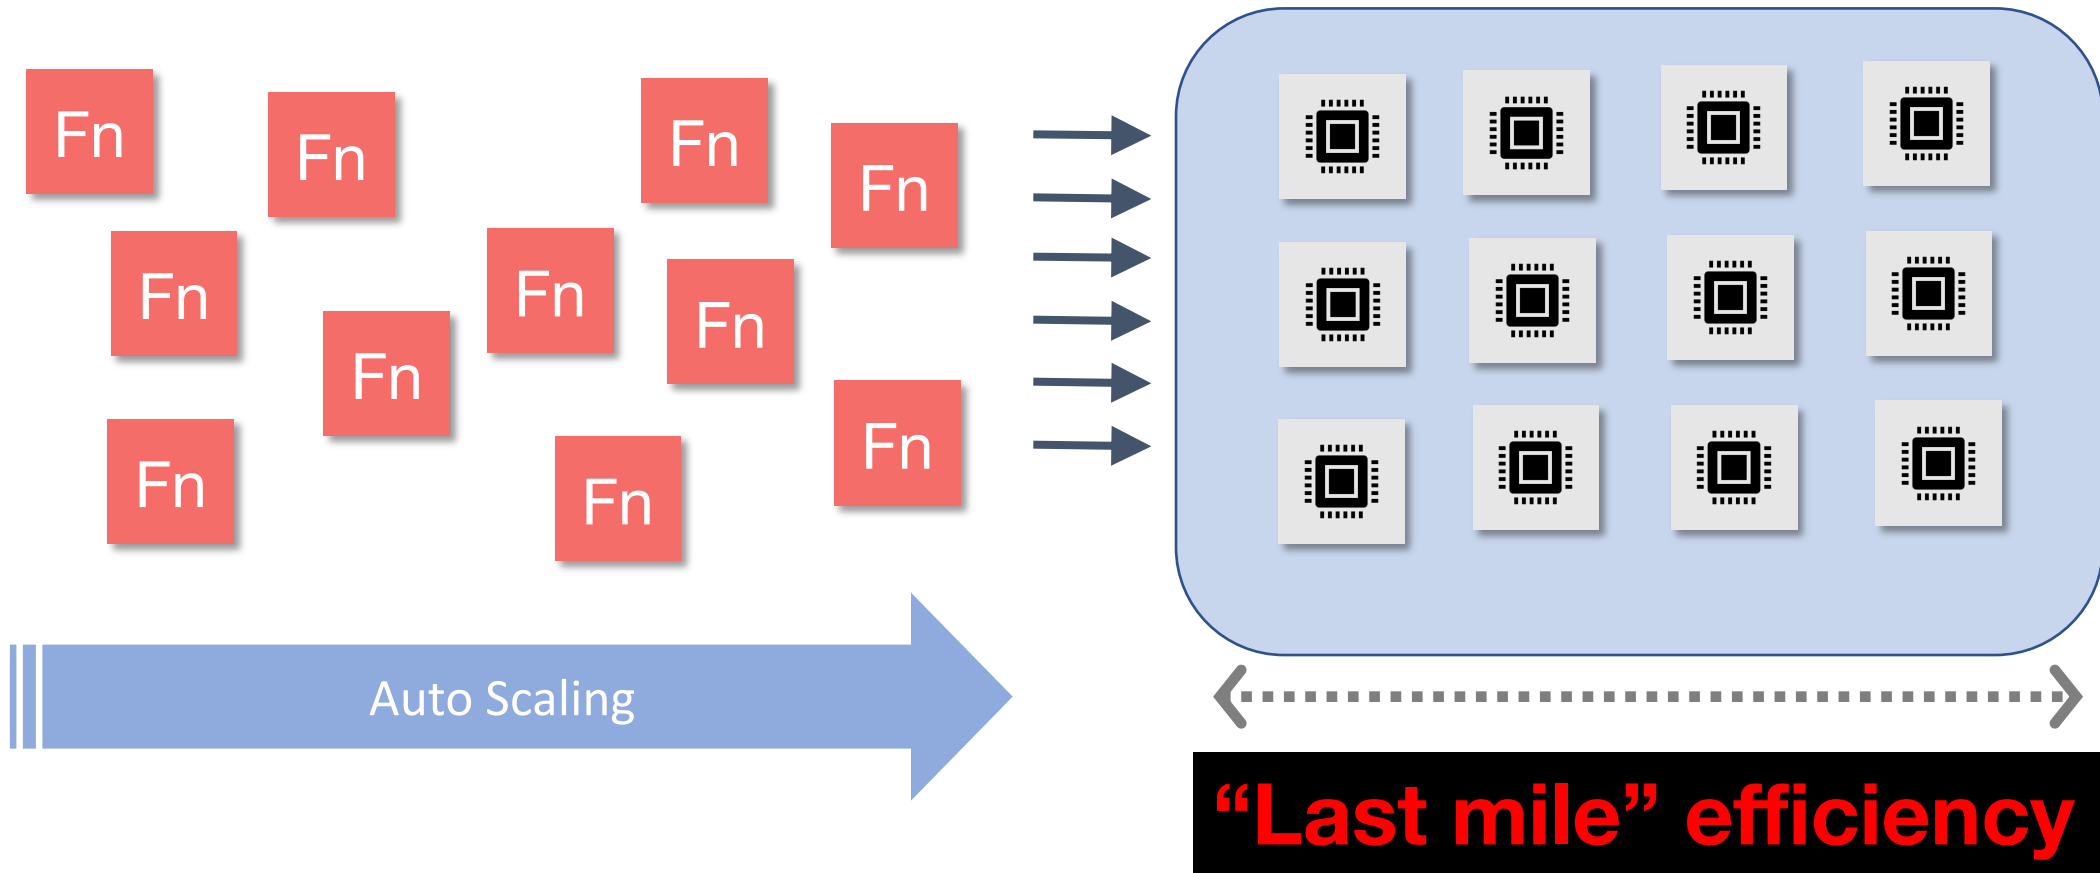
# Function-as-a-Service (FaaS)




# Function-as-a-Service (FaaS)



# Serverless functions eventually run in OSeS




# Production FaaS workloads are highly heterogeneous



A 14-day production FaaS workloads from Azure Function (ATC'20)

# Production FaaS workloads are highly heterogeneous

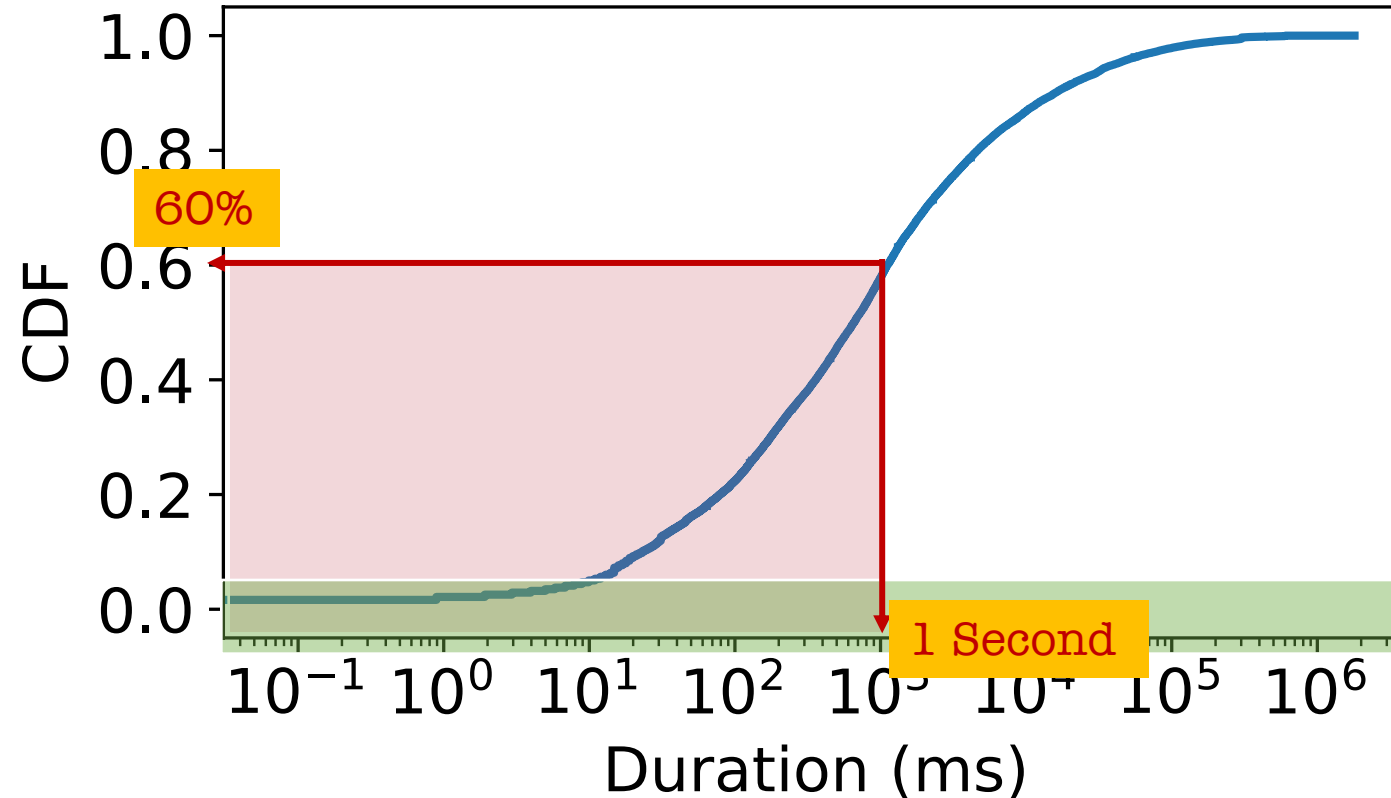


A 14-day production FaaS workloads from Azure Function (ATC'20)

- Features a mixture of short and long functions



# Production FaaS workloads are highly heterogeneous




A 14-day production FaaS workloads from Azure Function (ATC'20)

- Features a mixture of short and long functions
- A majority (60%) of functions finish in one second




# Poor performance under existing OS scheduling



## CFS

- \* Proportional-Share time slice
- \* Default Linux Scheduler

# Poor performance under existing OS scheduling




## SRTF

\* Optimal

\* Not practical

# Poor performance under existing OS scheduling




## IDEAL

\* Infinite resource

\* No preemption




# Poor performance under existing OS scheduling




Proportional-Share Scheduling  
offer poor performance

# Poor performance under existing OS scheduling



Proportional-Share Scheduling  
offer poor performance



**Function Run-Time Effectiveness**  
= Service time / Turnaround time

CFS frequently preempts functions,  
causing longer waiting time (w/  
smaller RTEs)



Poor performance under existing OS scheduling

## **Existing Linux scheduling policies are a poor match for emerging FaaS workloads**

- **Implication #1:** OS-level function scheduling must be workload-aware
- **Implication #2:** Approximating SRTF (shortest remaining time first) will provide a significant performance boost for short functions

# Outline

- Design
- Evaluation
- Conclusion







# Smart function scheduler (SFS)

SFS is a **FaaS-aware**, user-space OS scheduler

**Key idea:** Two-level scheduling



**FILTER**


1<sup>st</sup> Level Scheduling



# Smart function scheduler (SFS)

SFS is a **FaaS-aware**, user-space OS scheduler


**Key idea:** Two-level scheduling



# Smart function scheduler (SFS)

SFS is a **FaaS-aware**, user-space OS scheduler

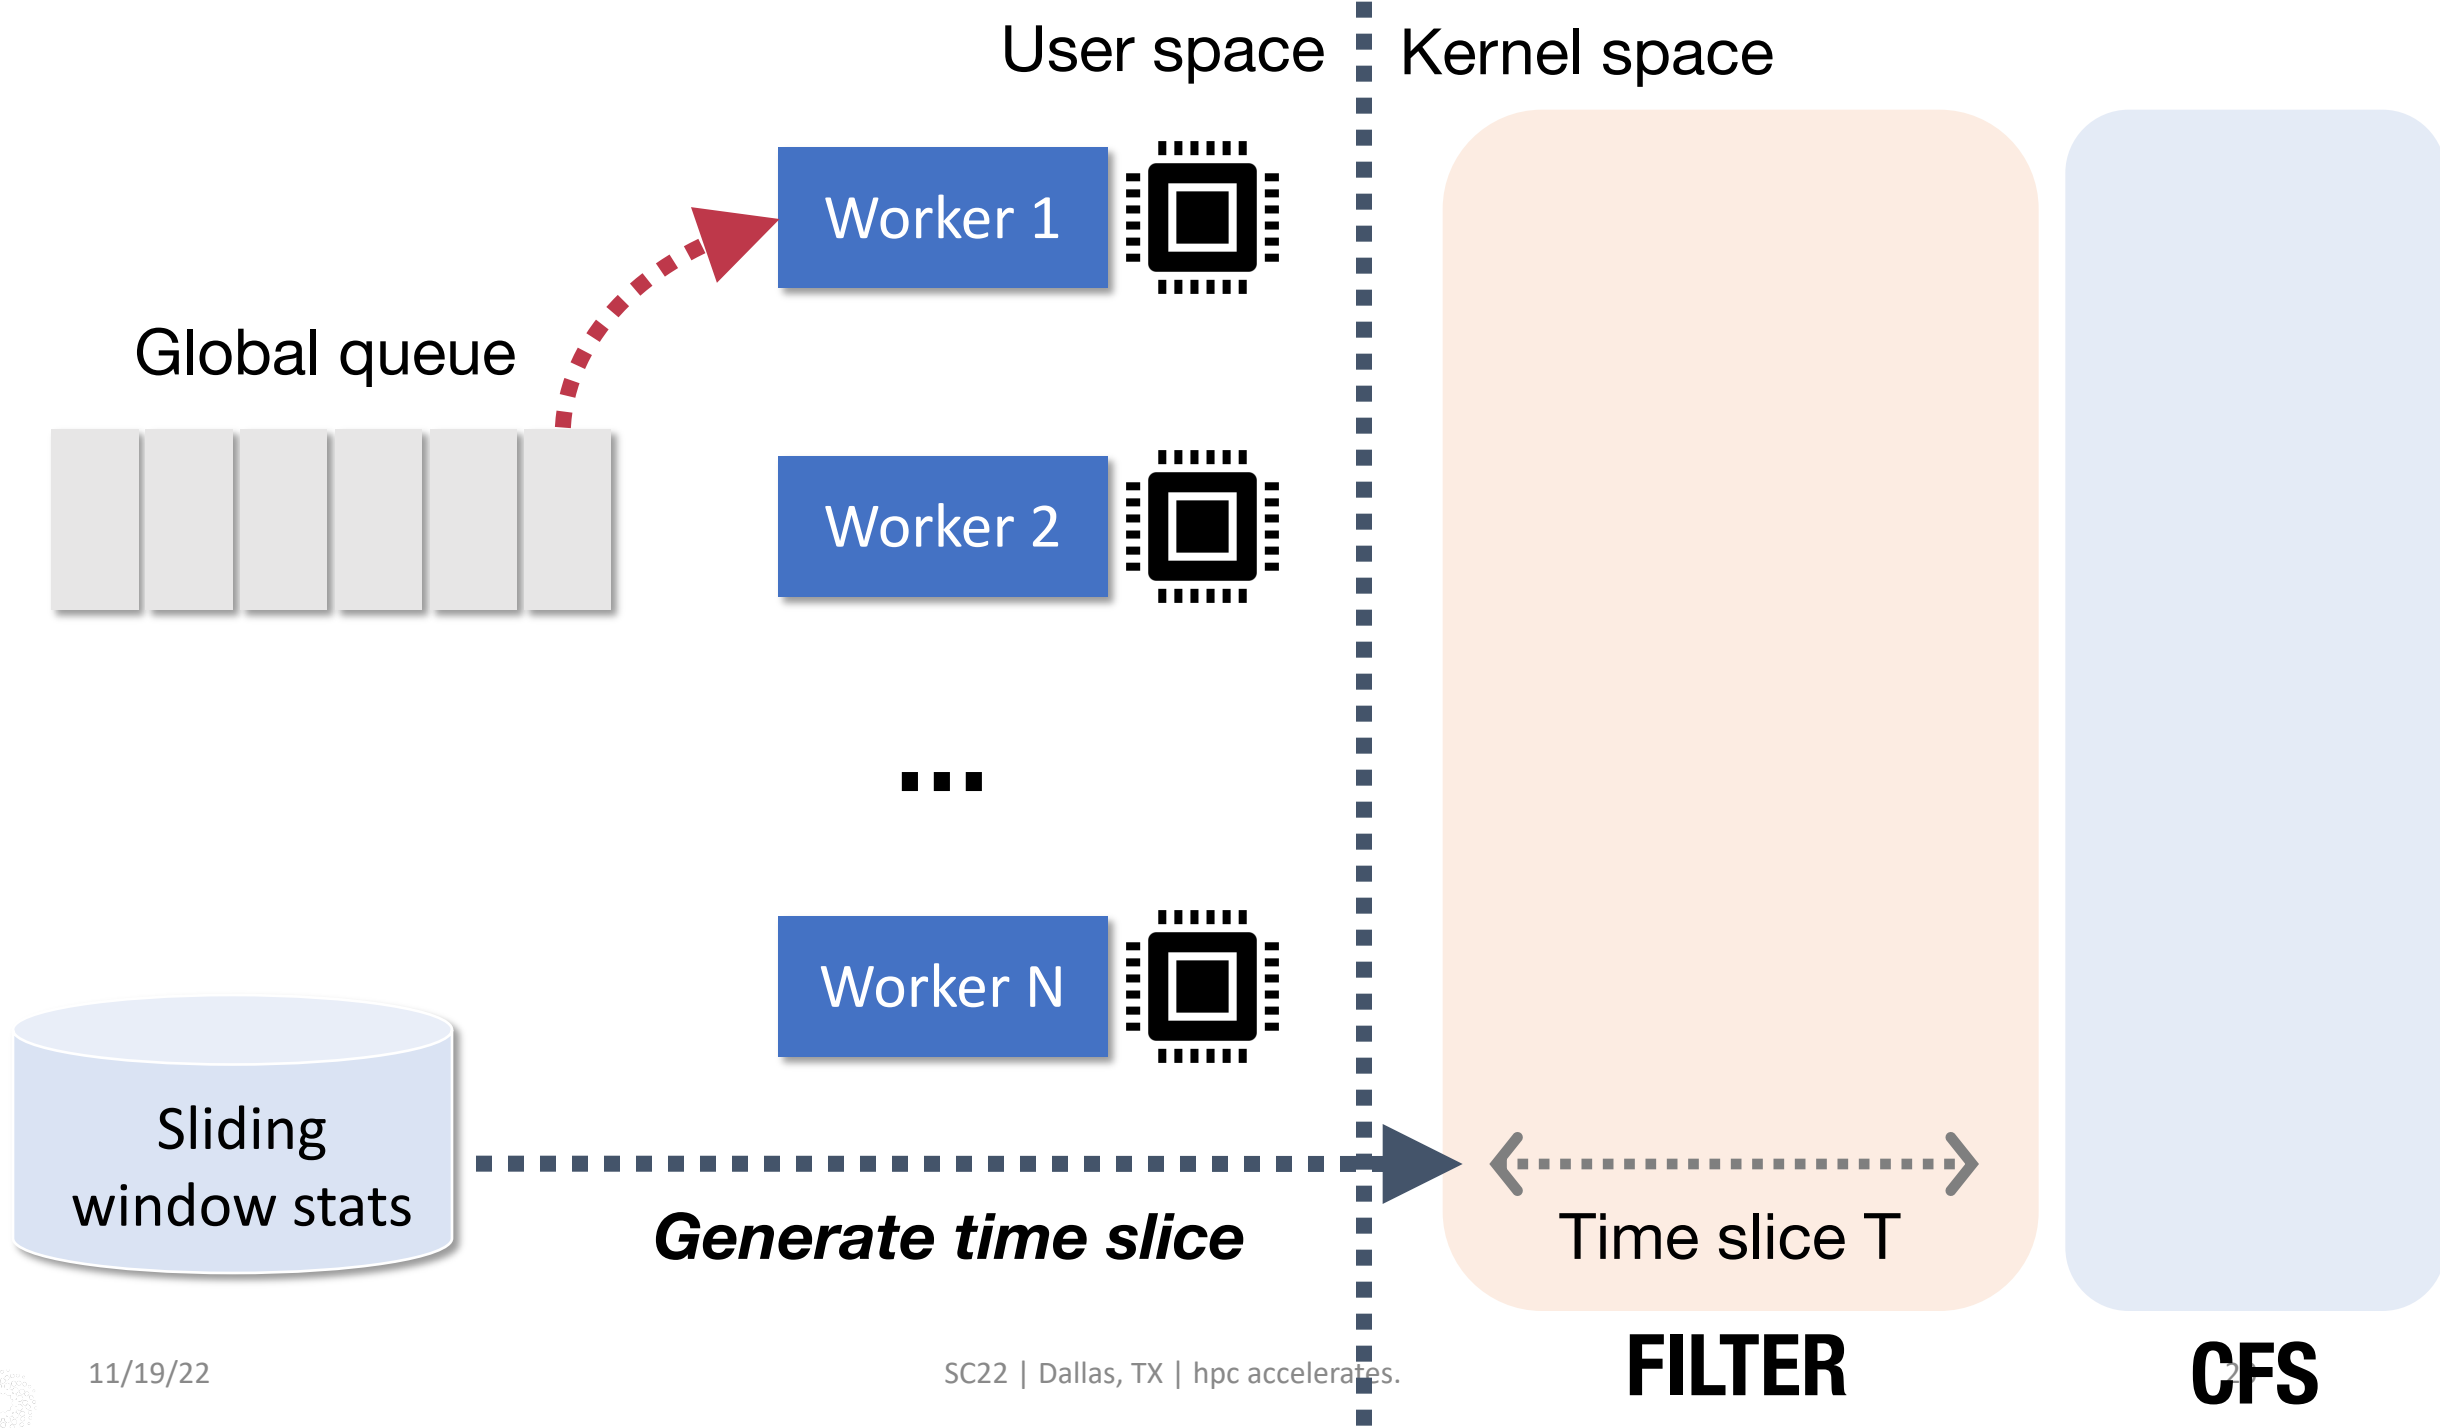
**Key idea:** Two-level scheduling

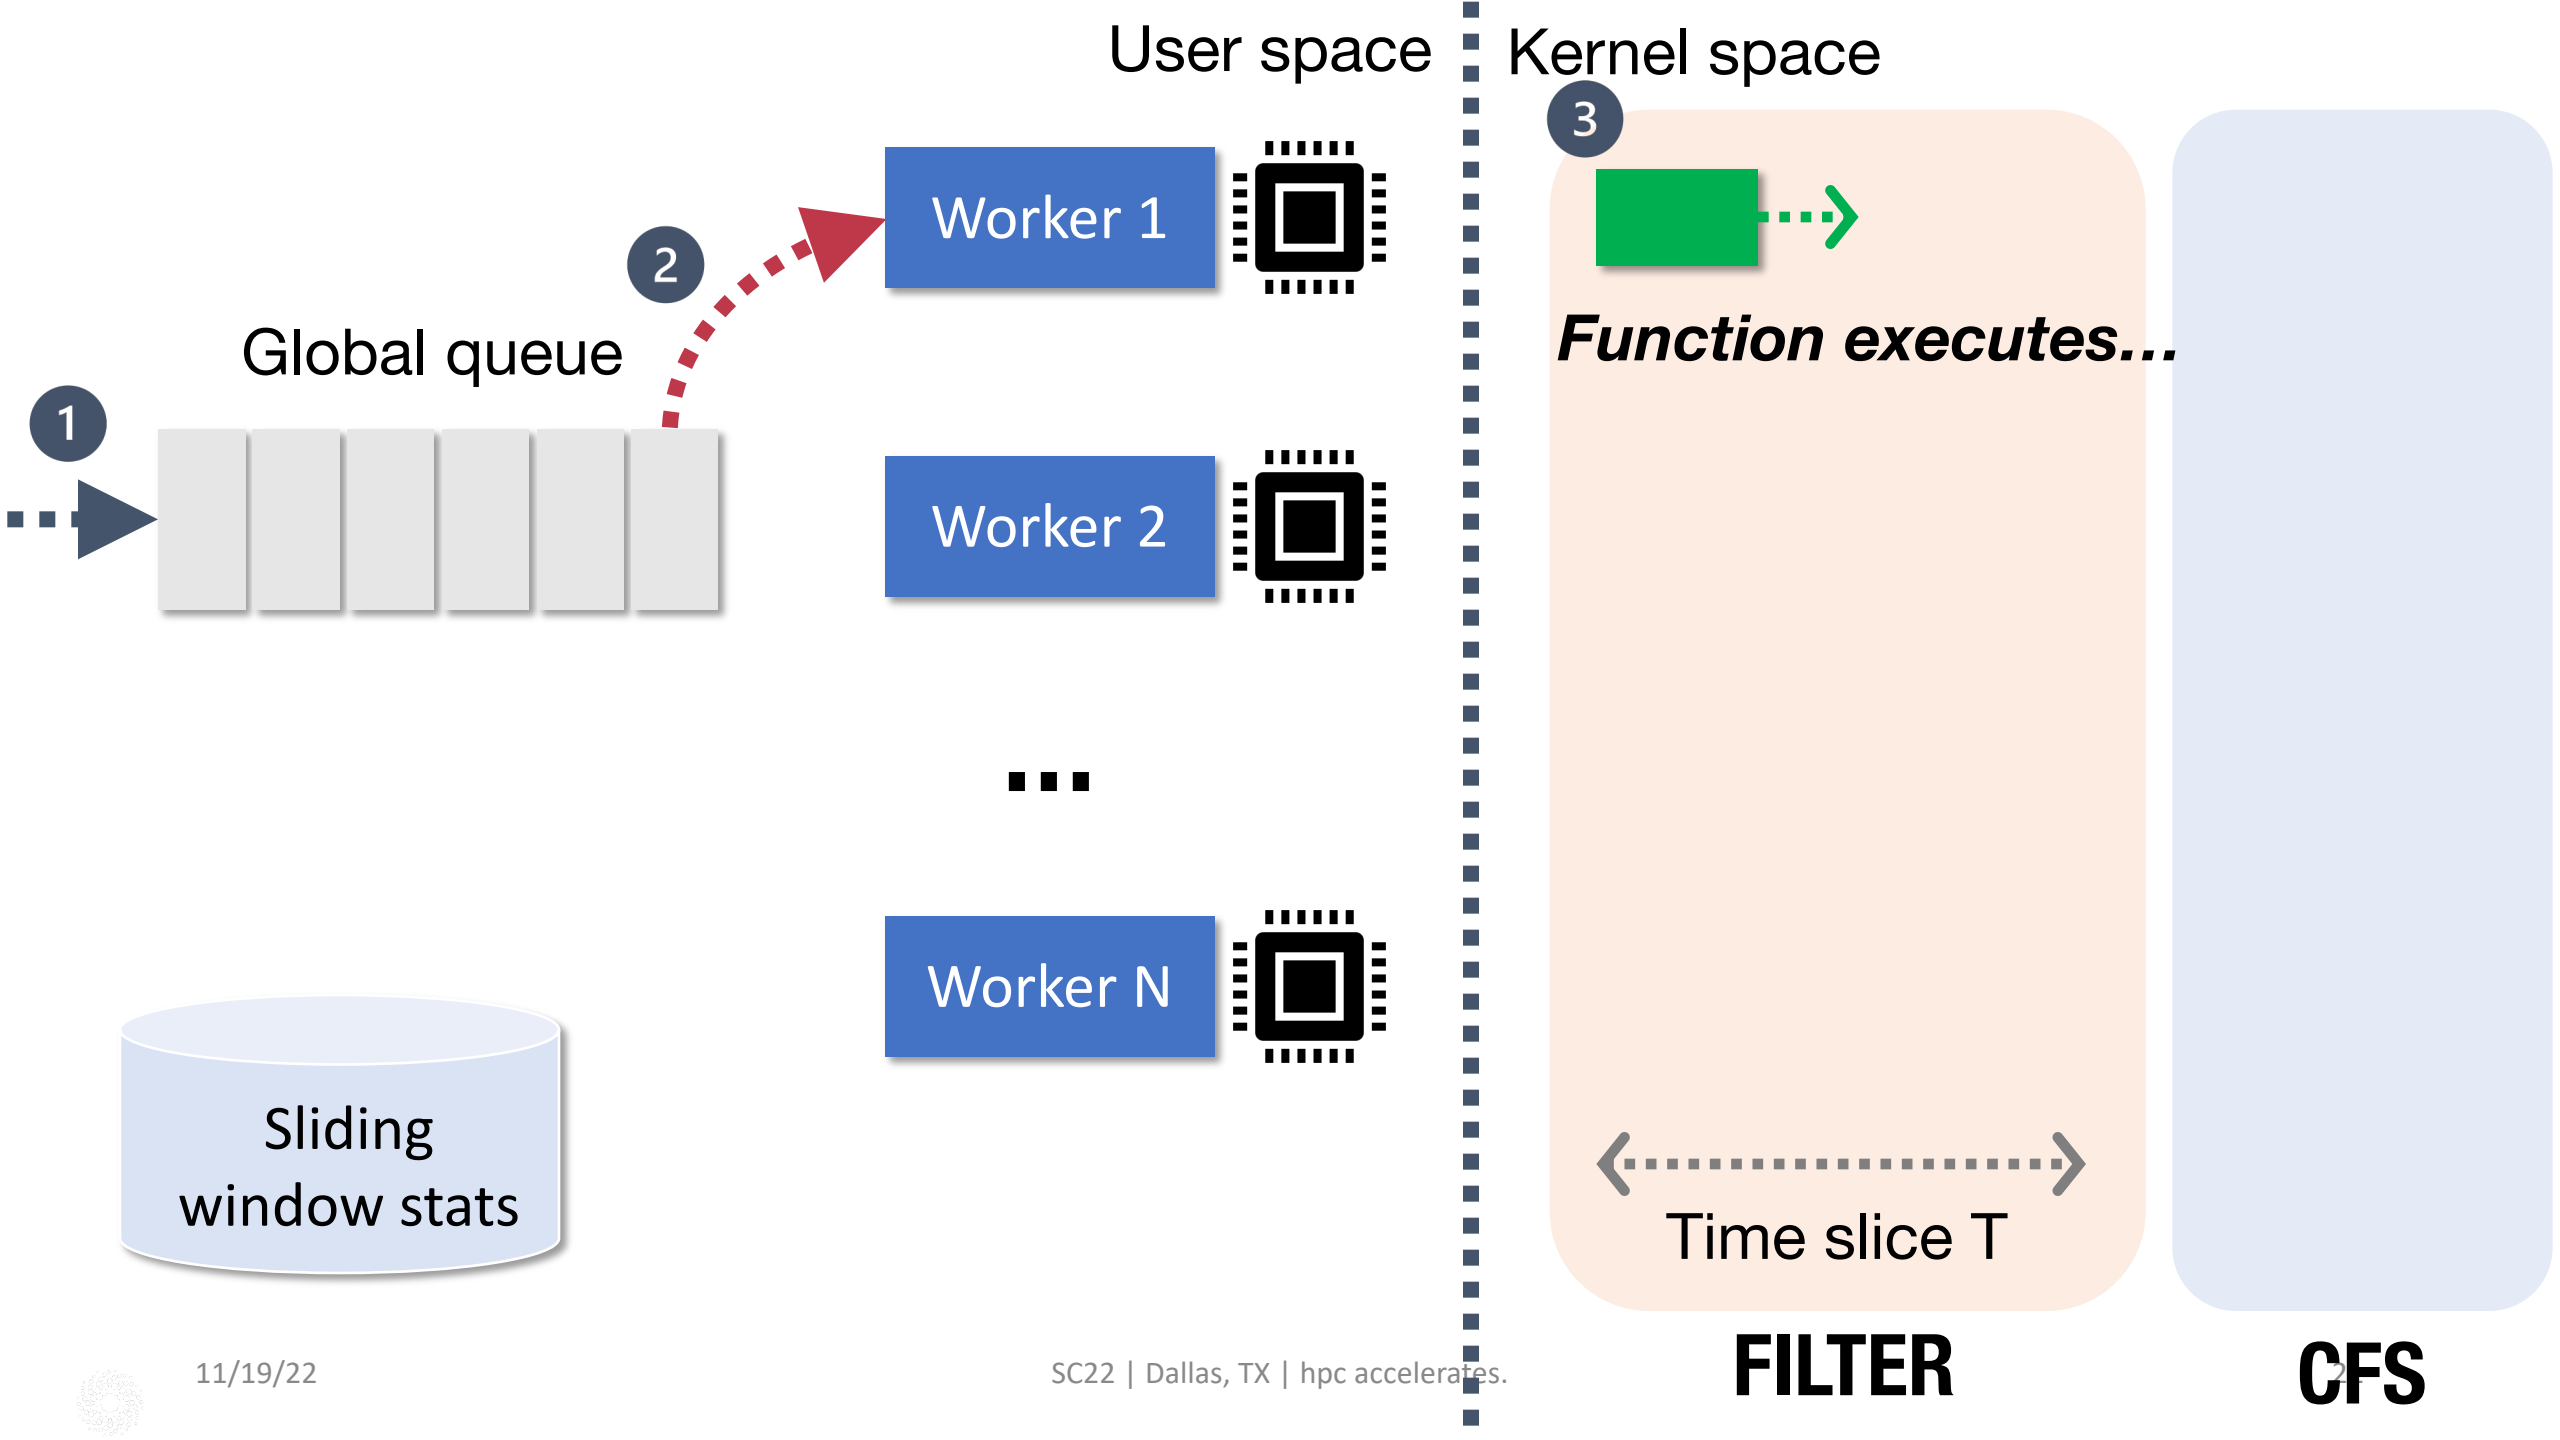


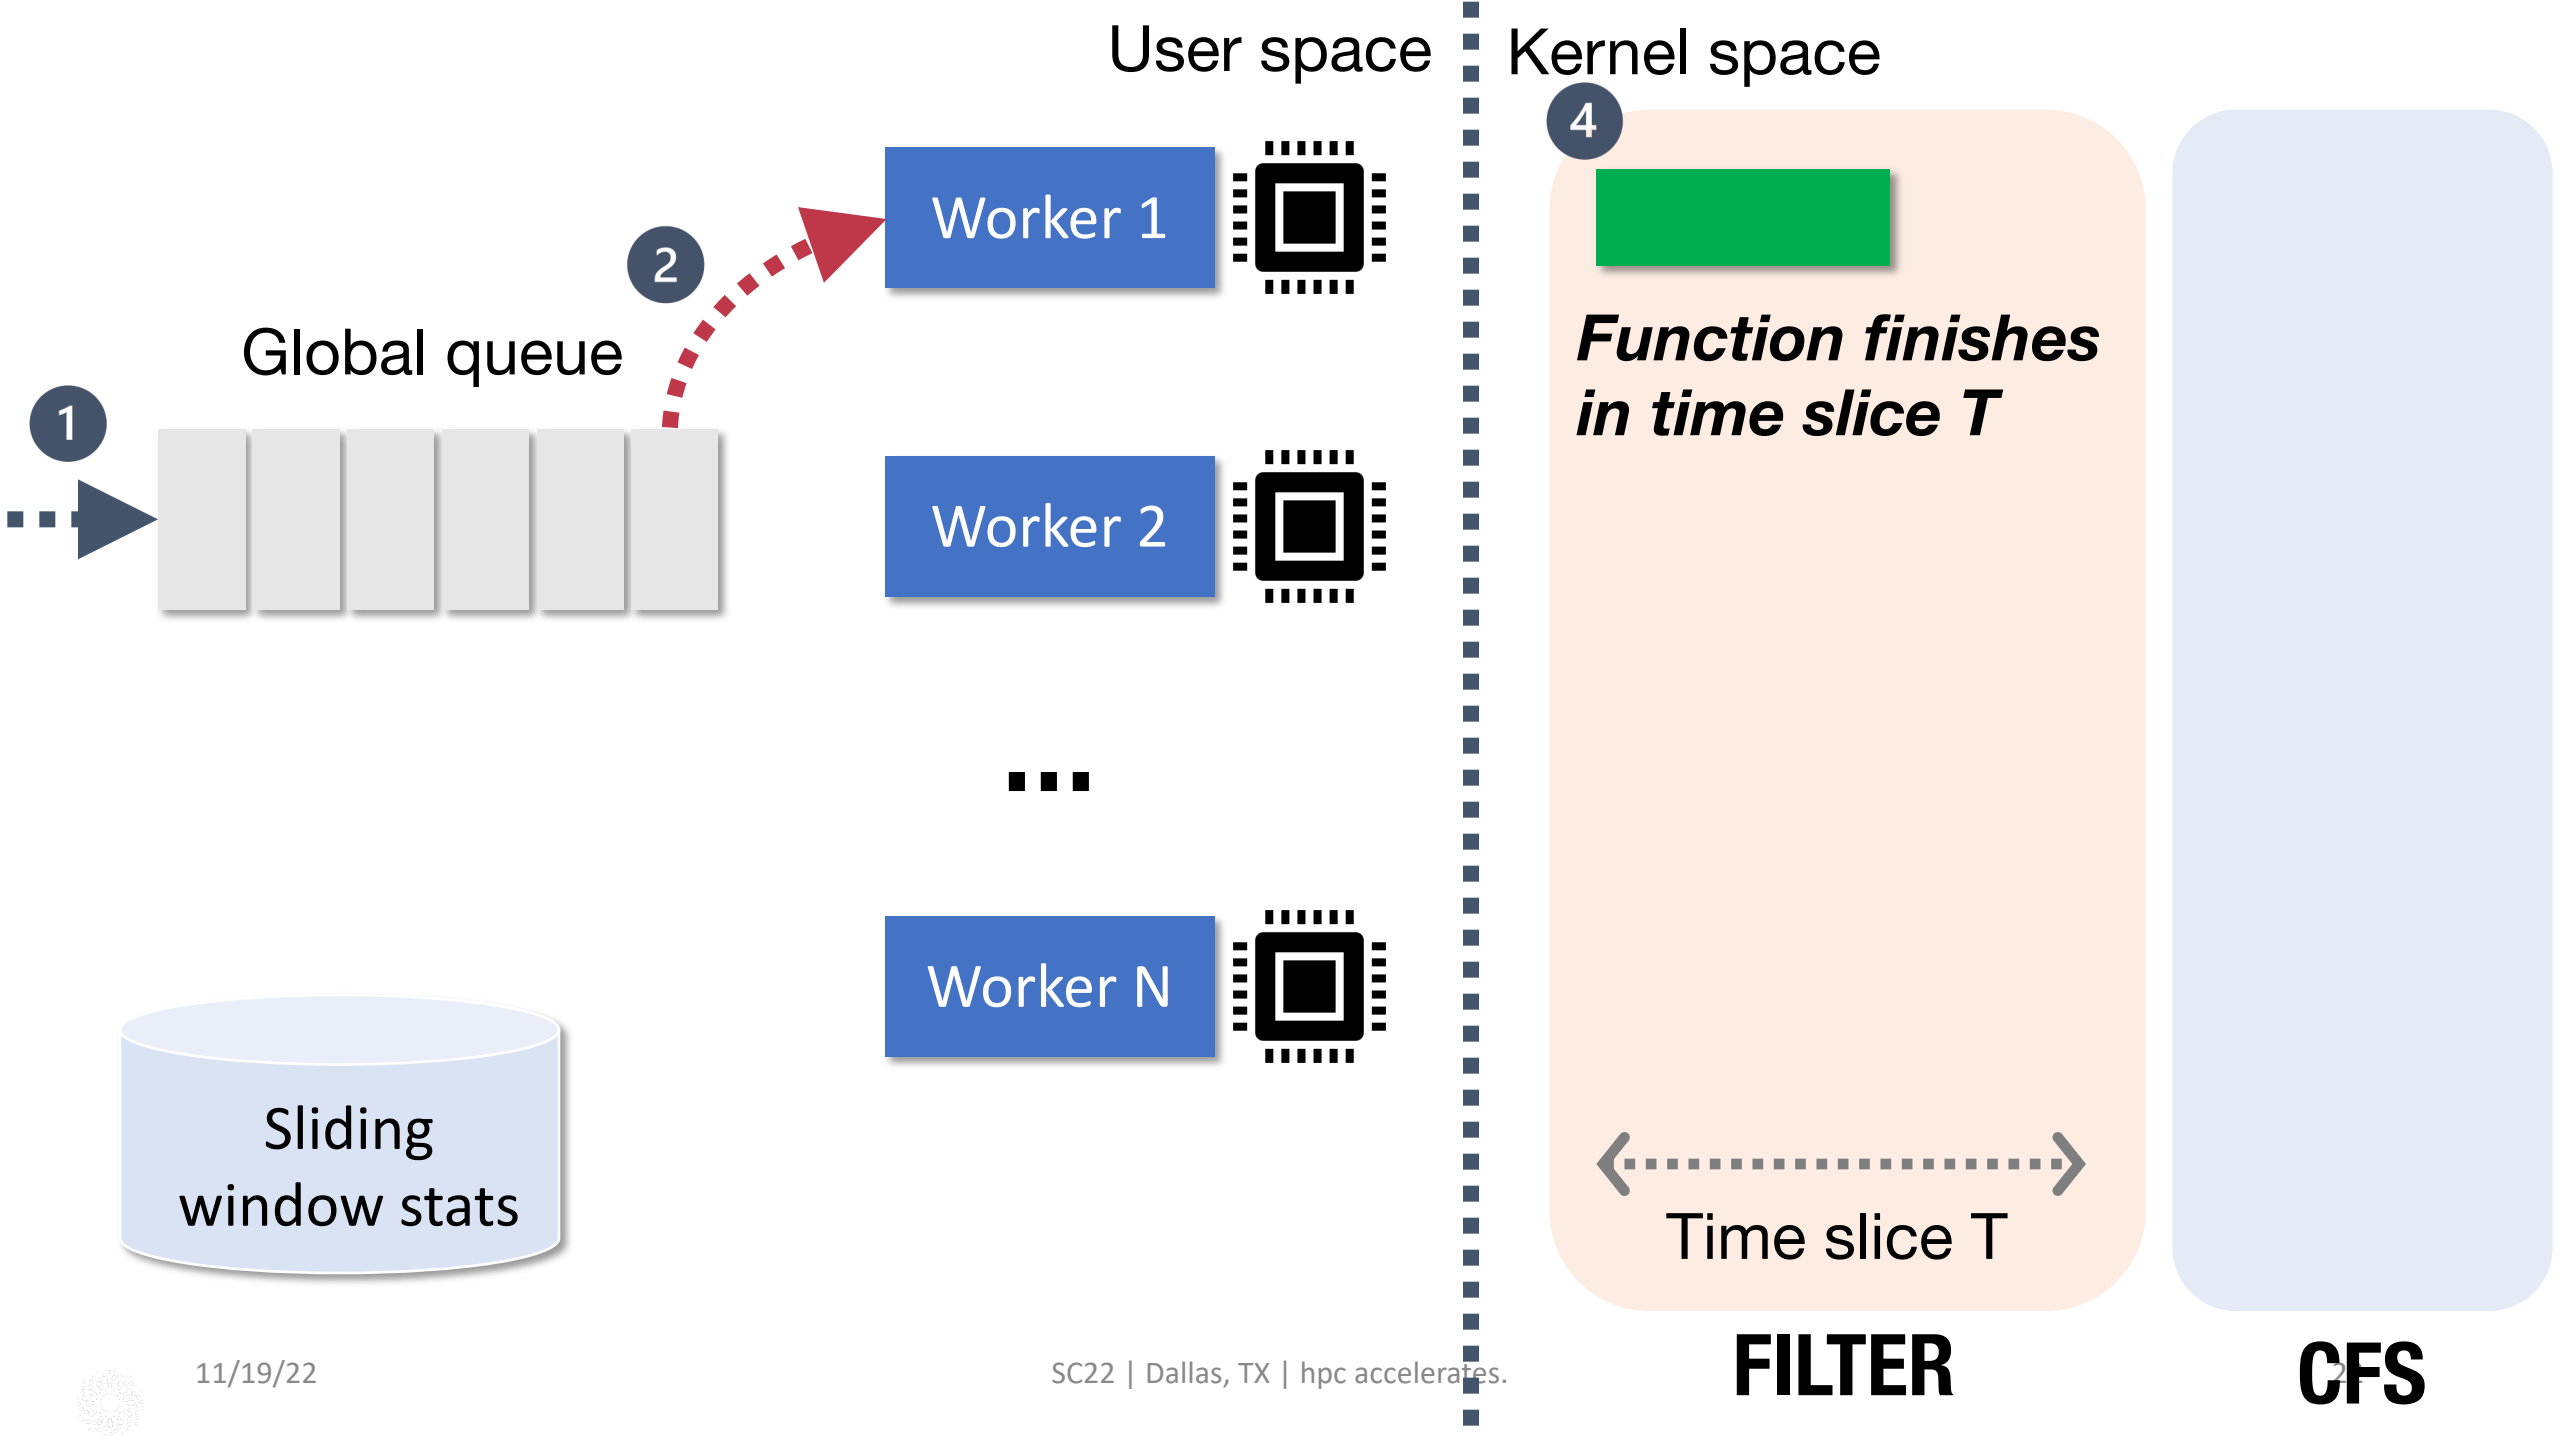
2<sup>nd</sup> Level Scheduling (CFS)

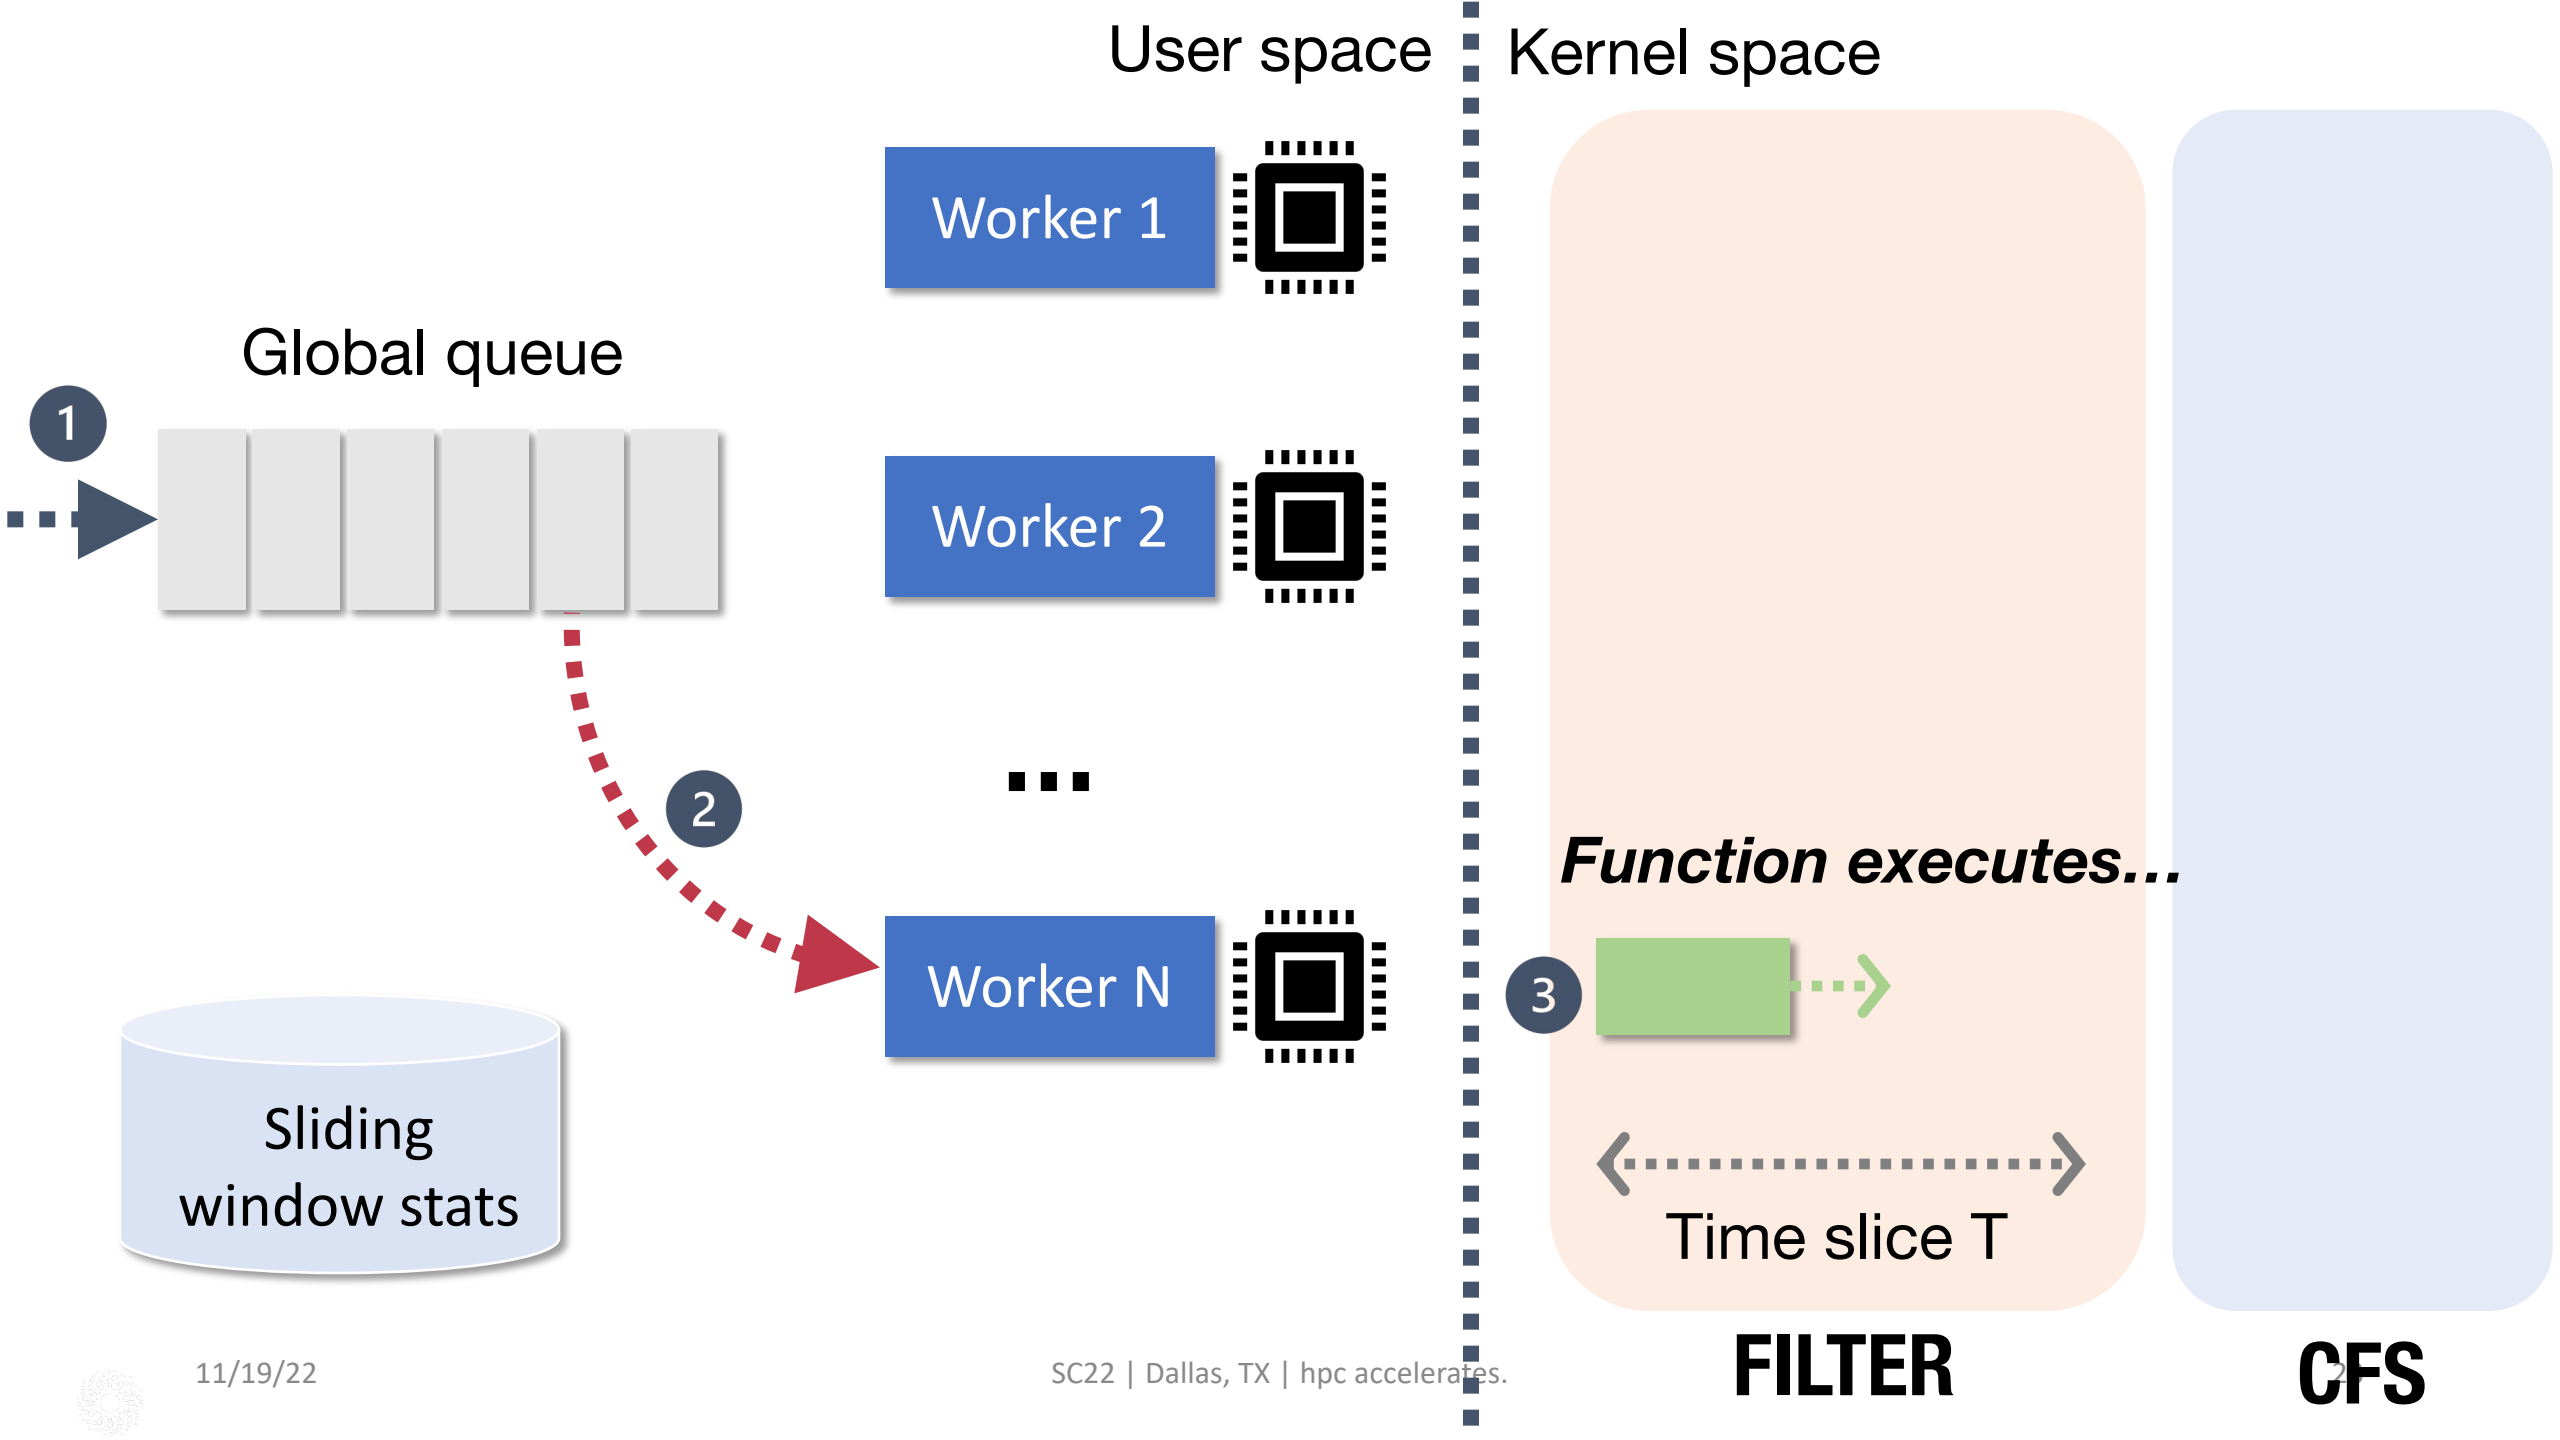
SC22 | Dallas, TX | hpc accelerates.

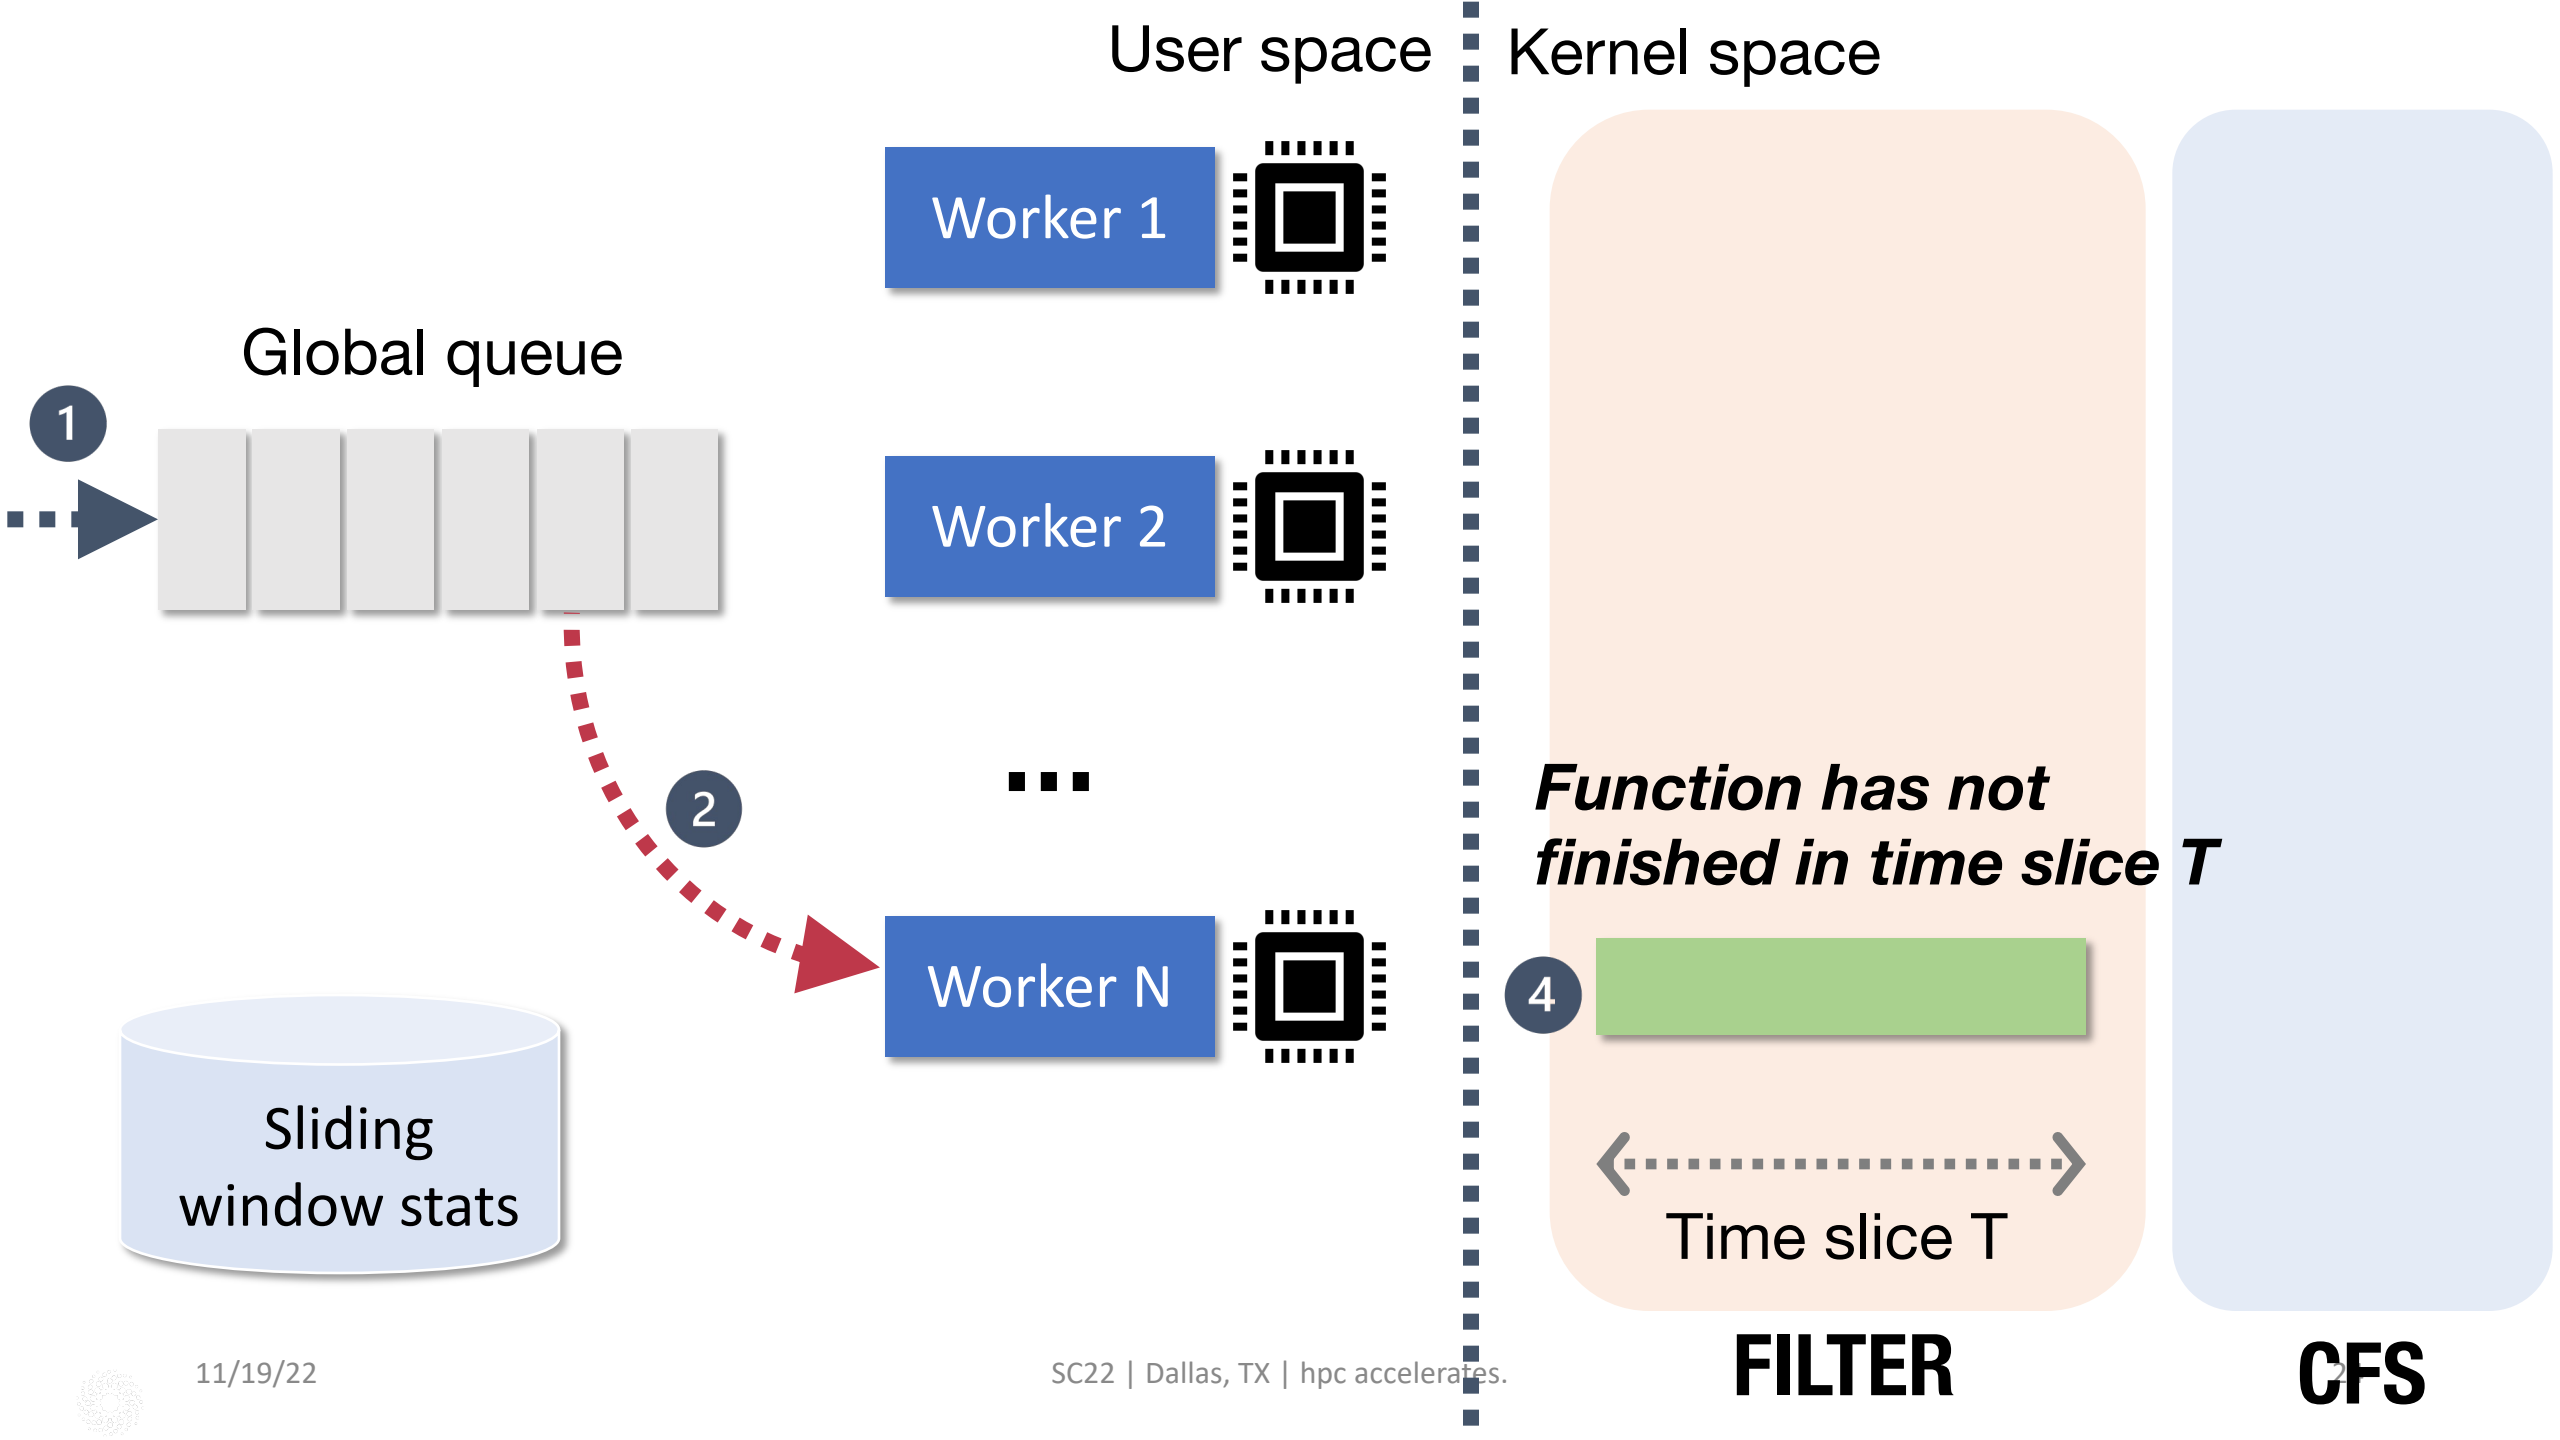




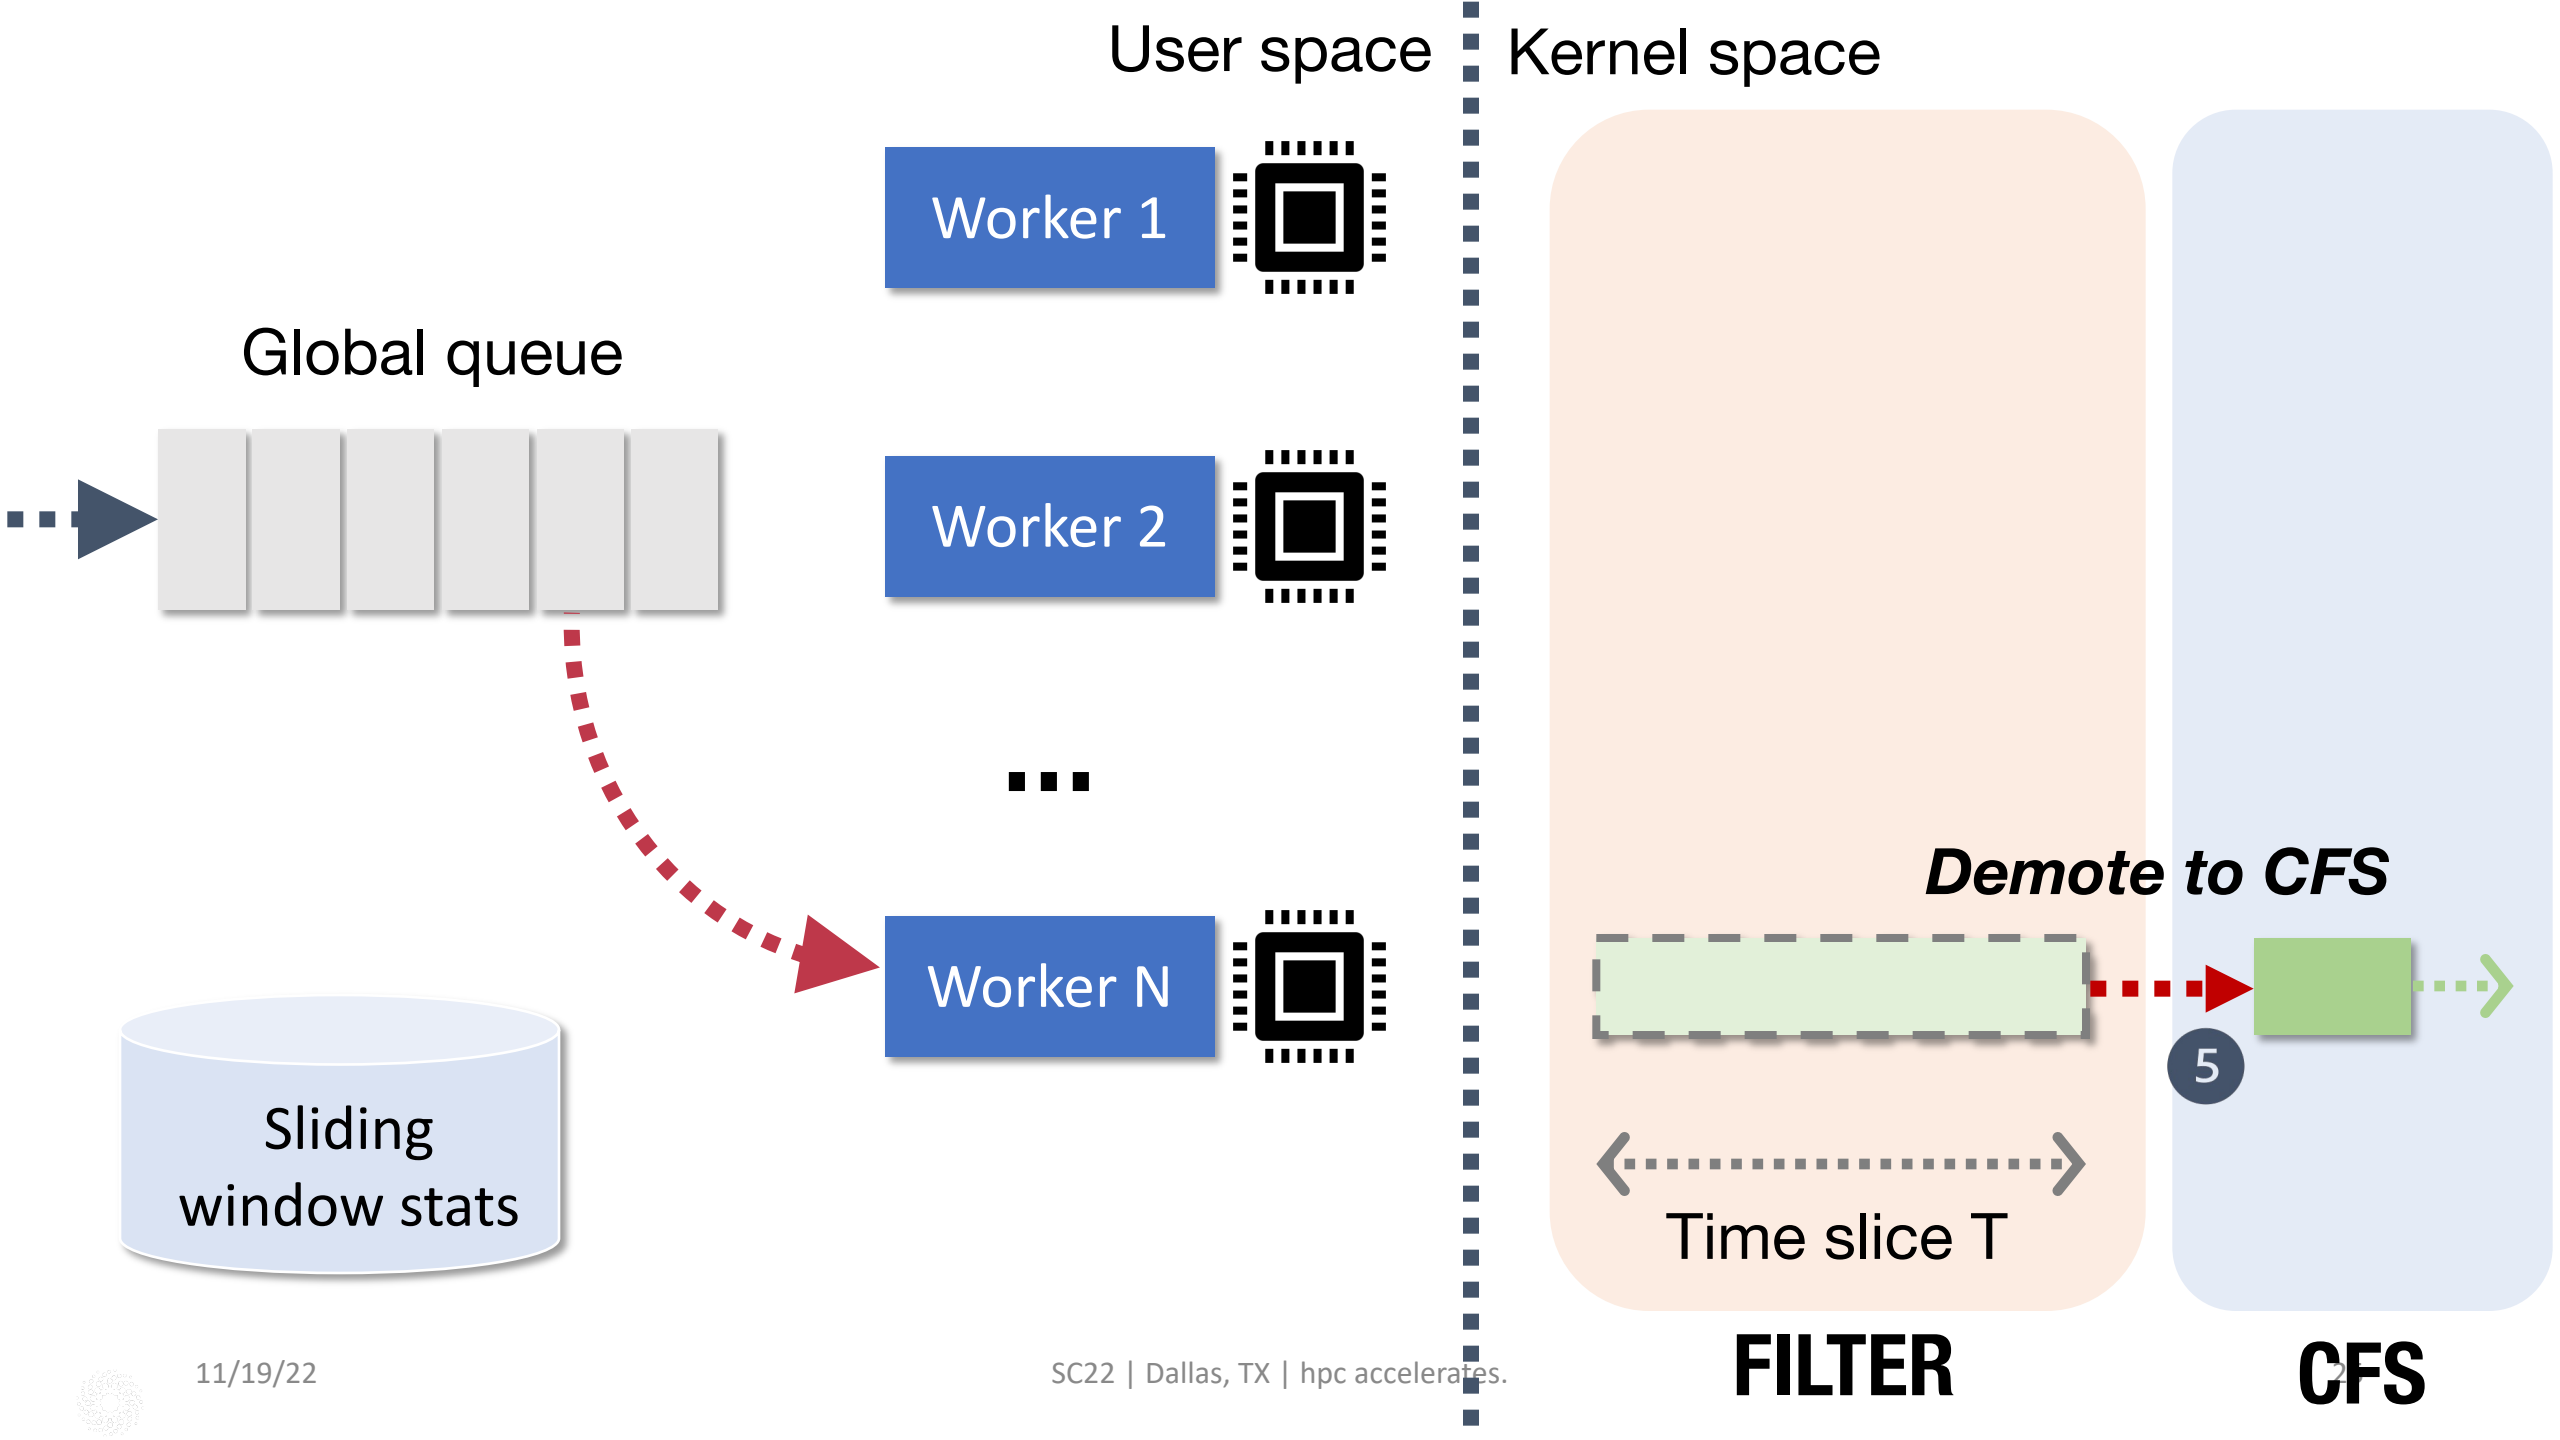


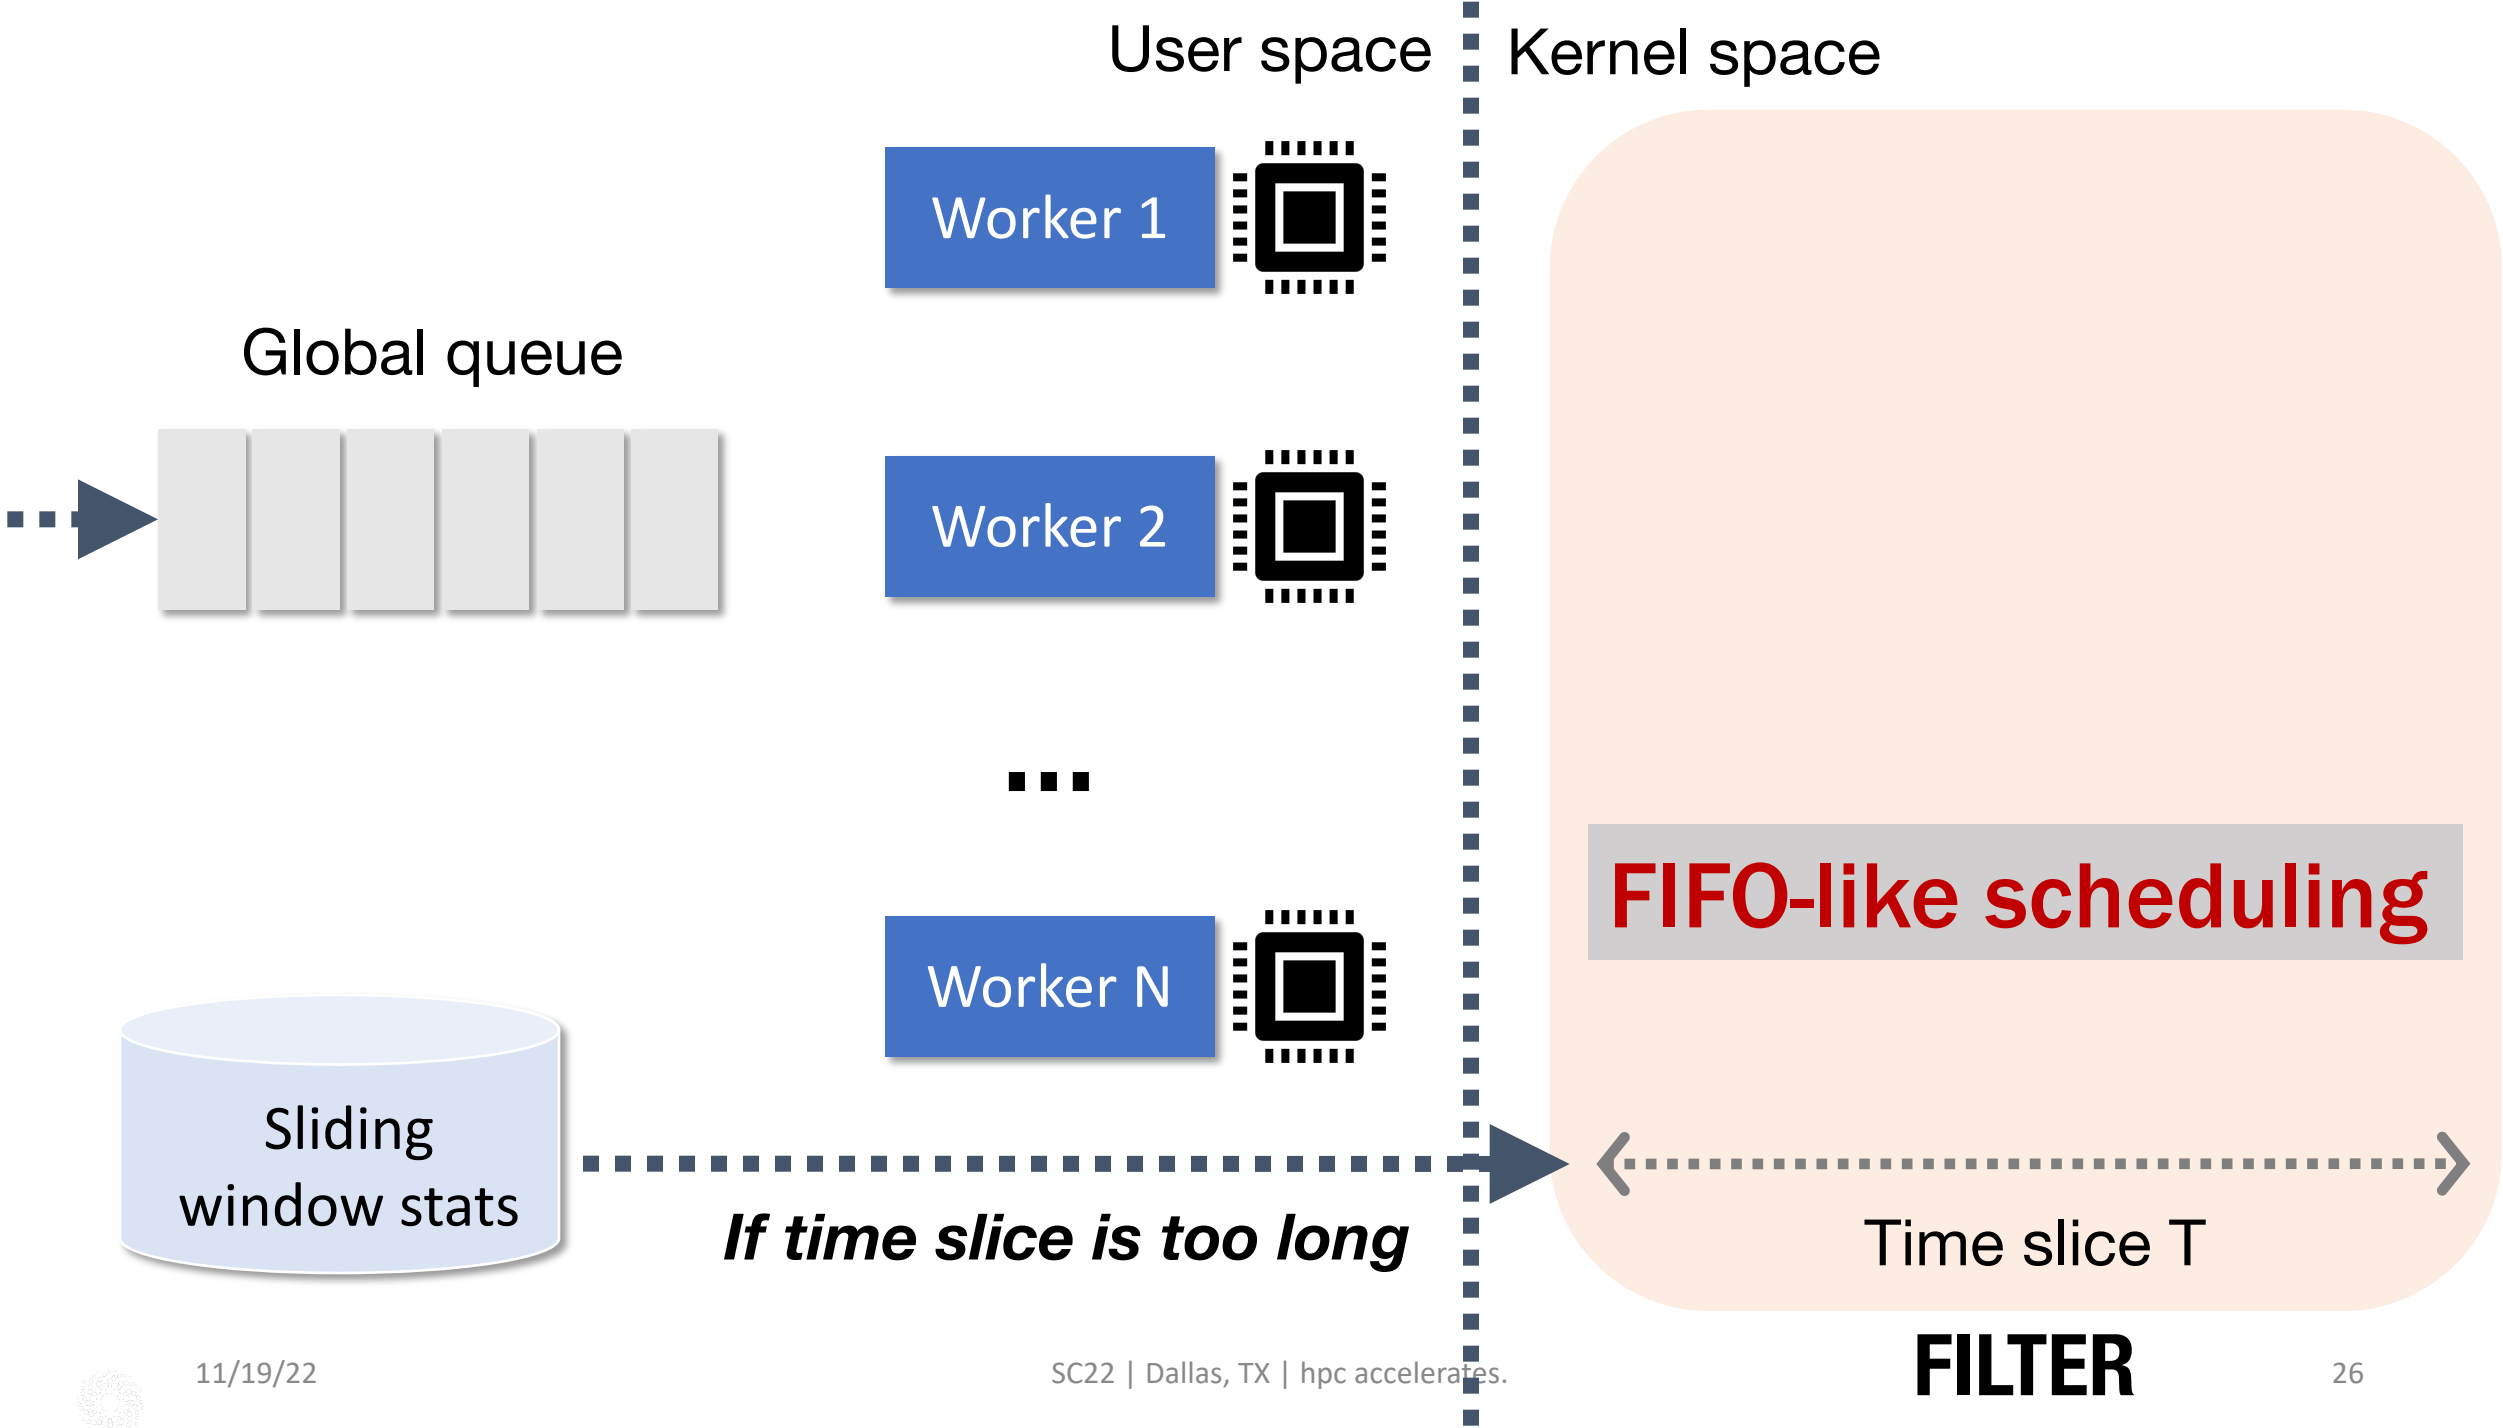


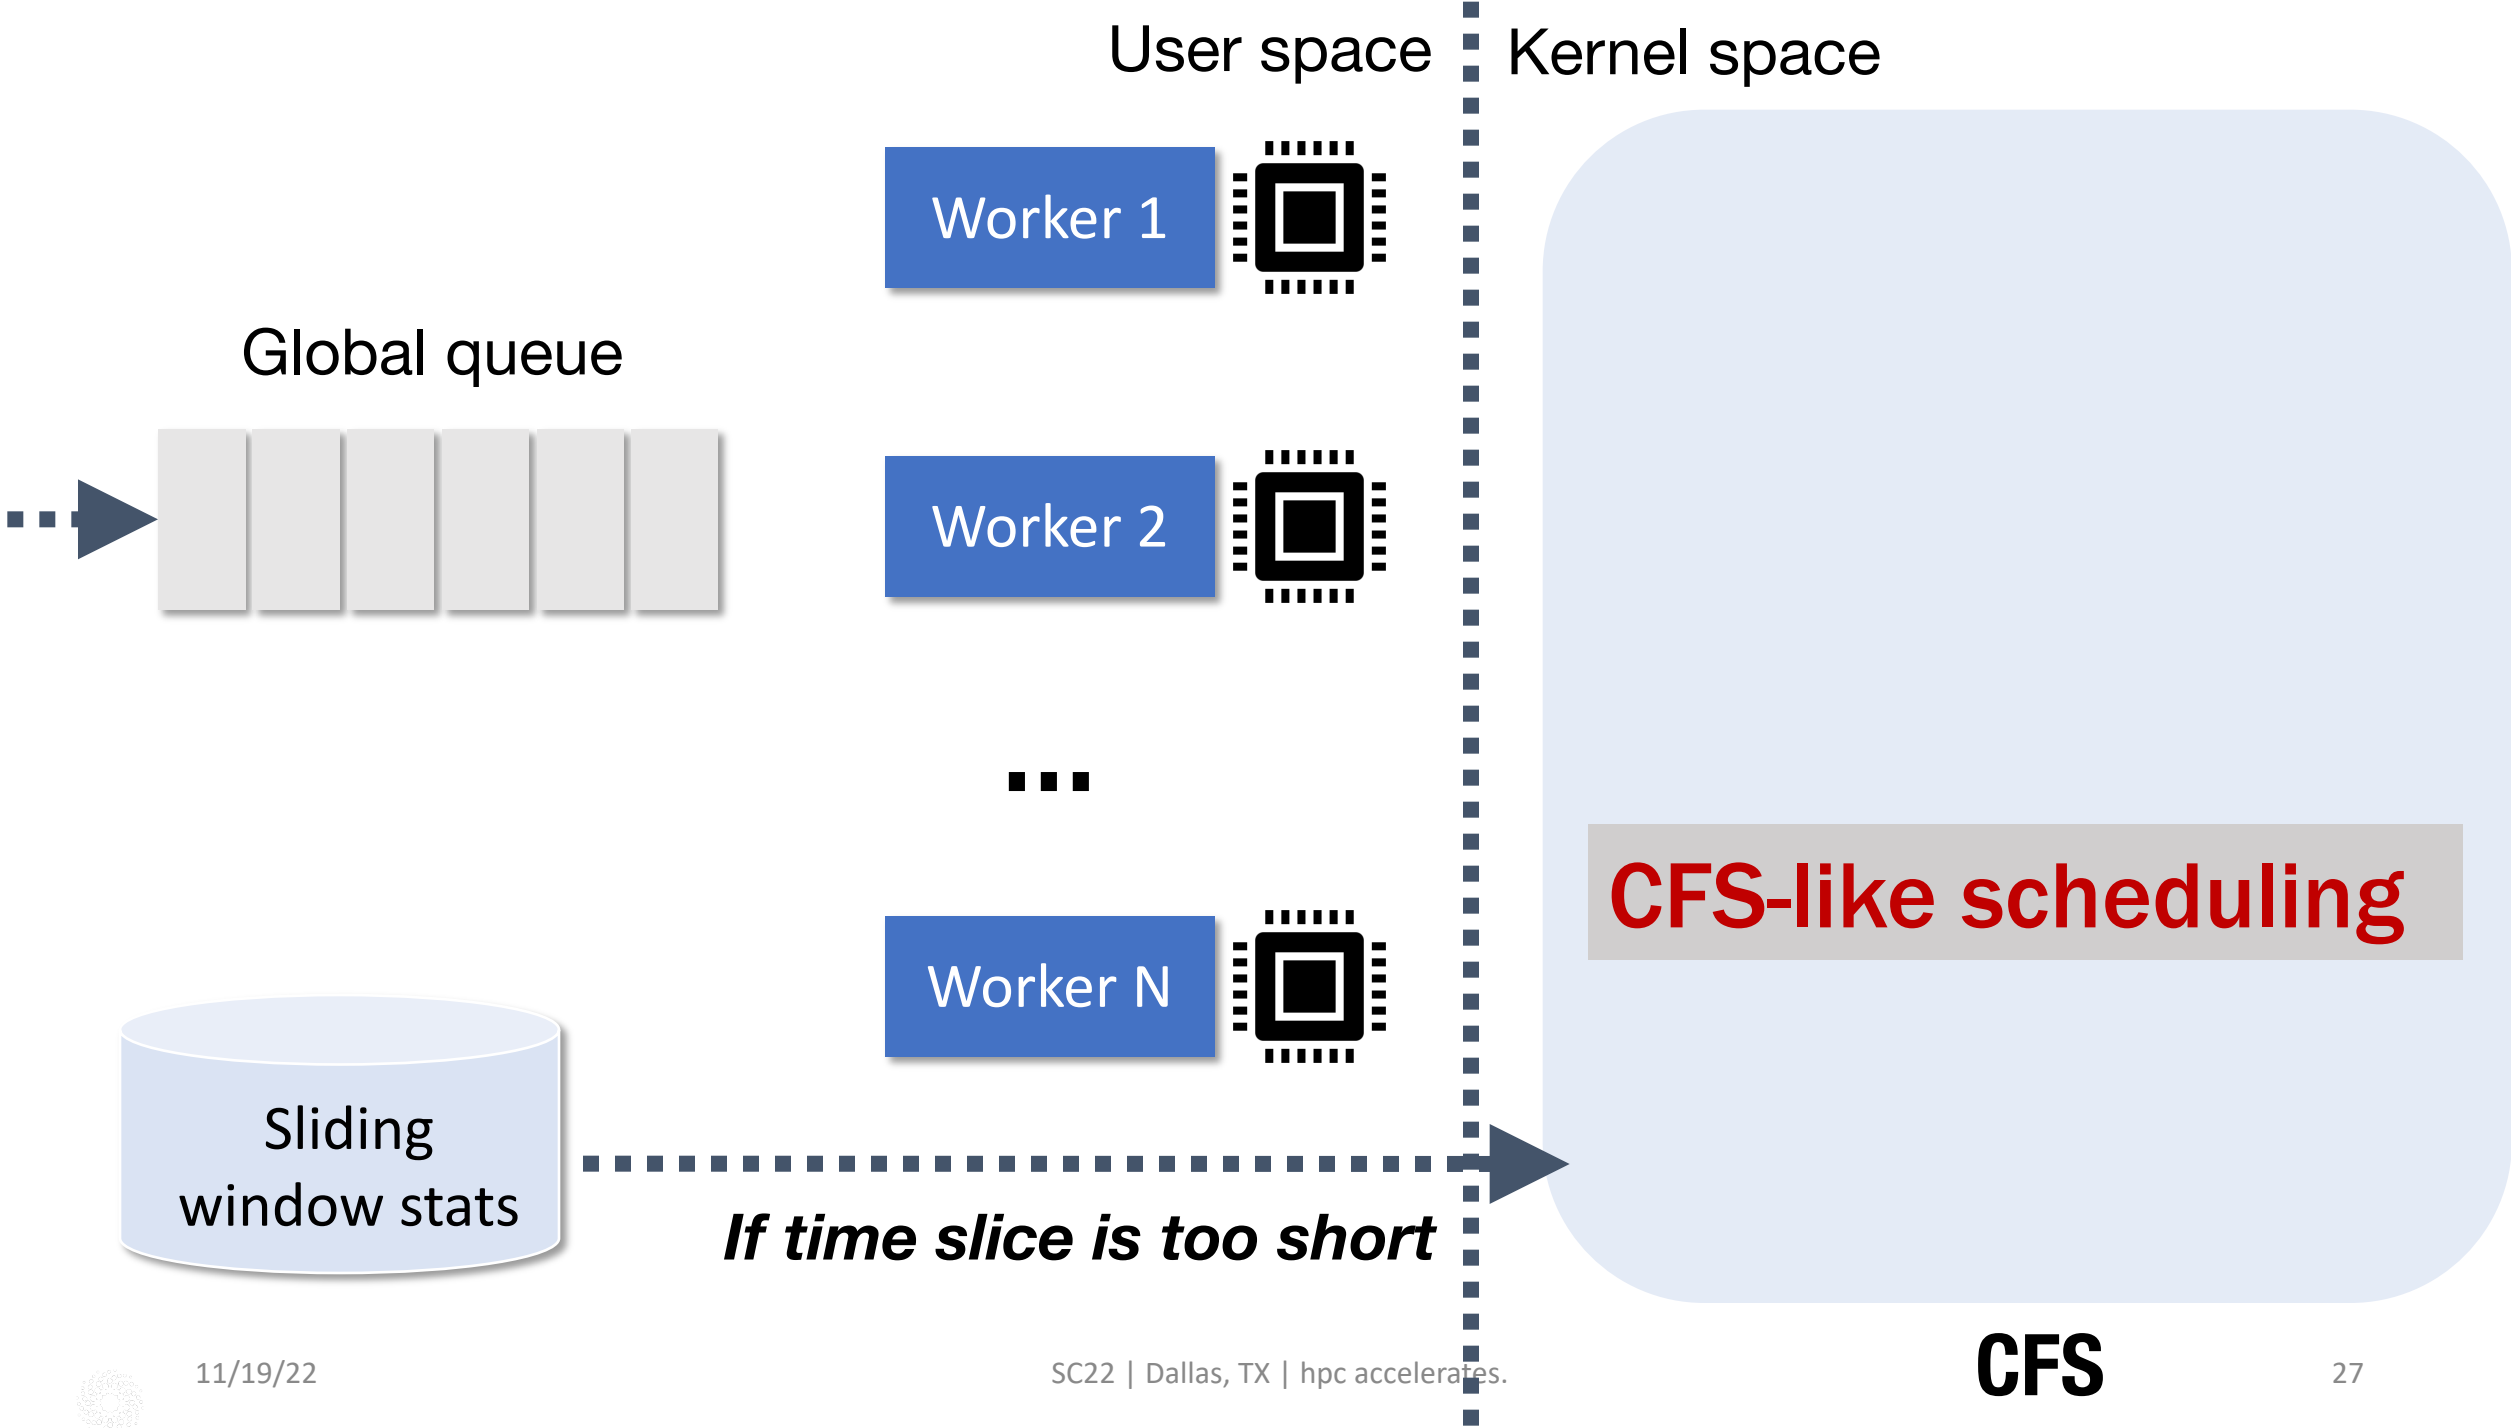


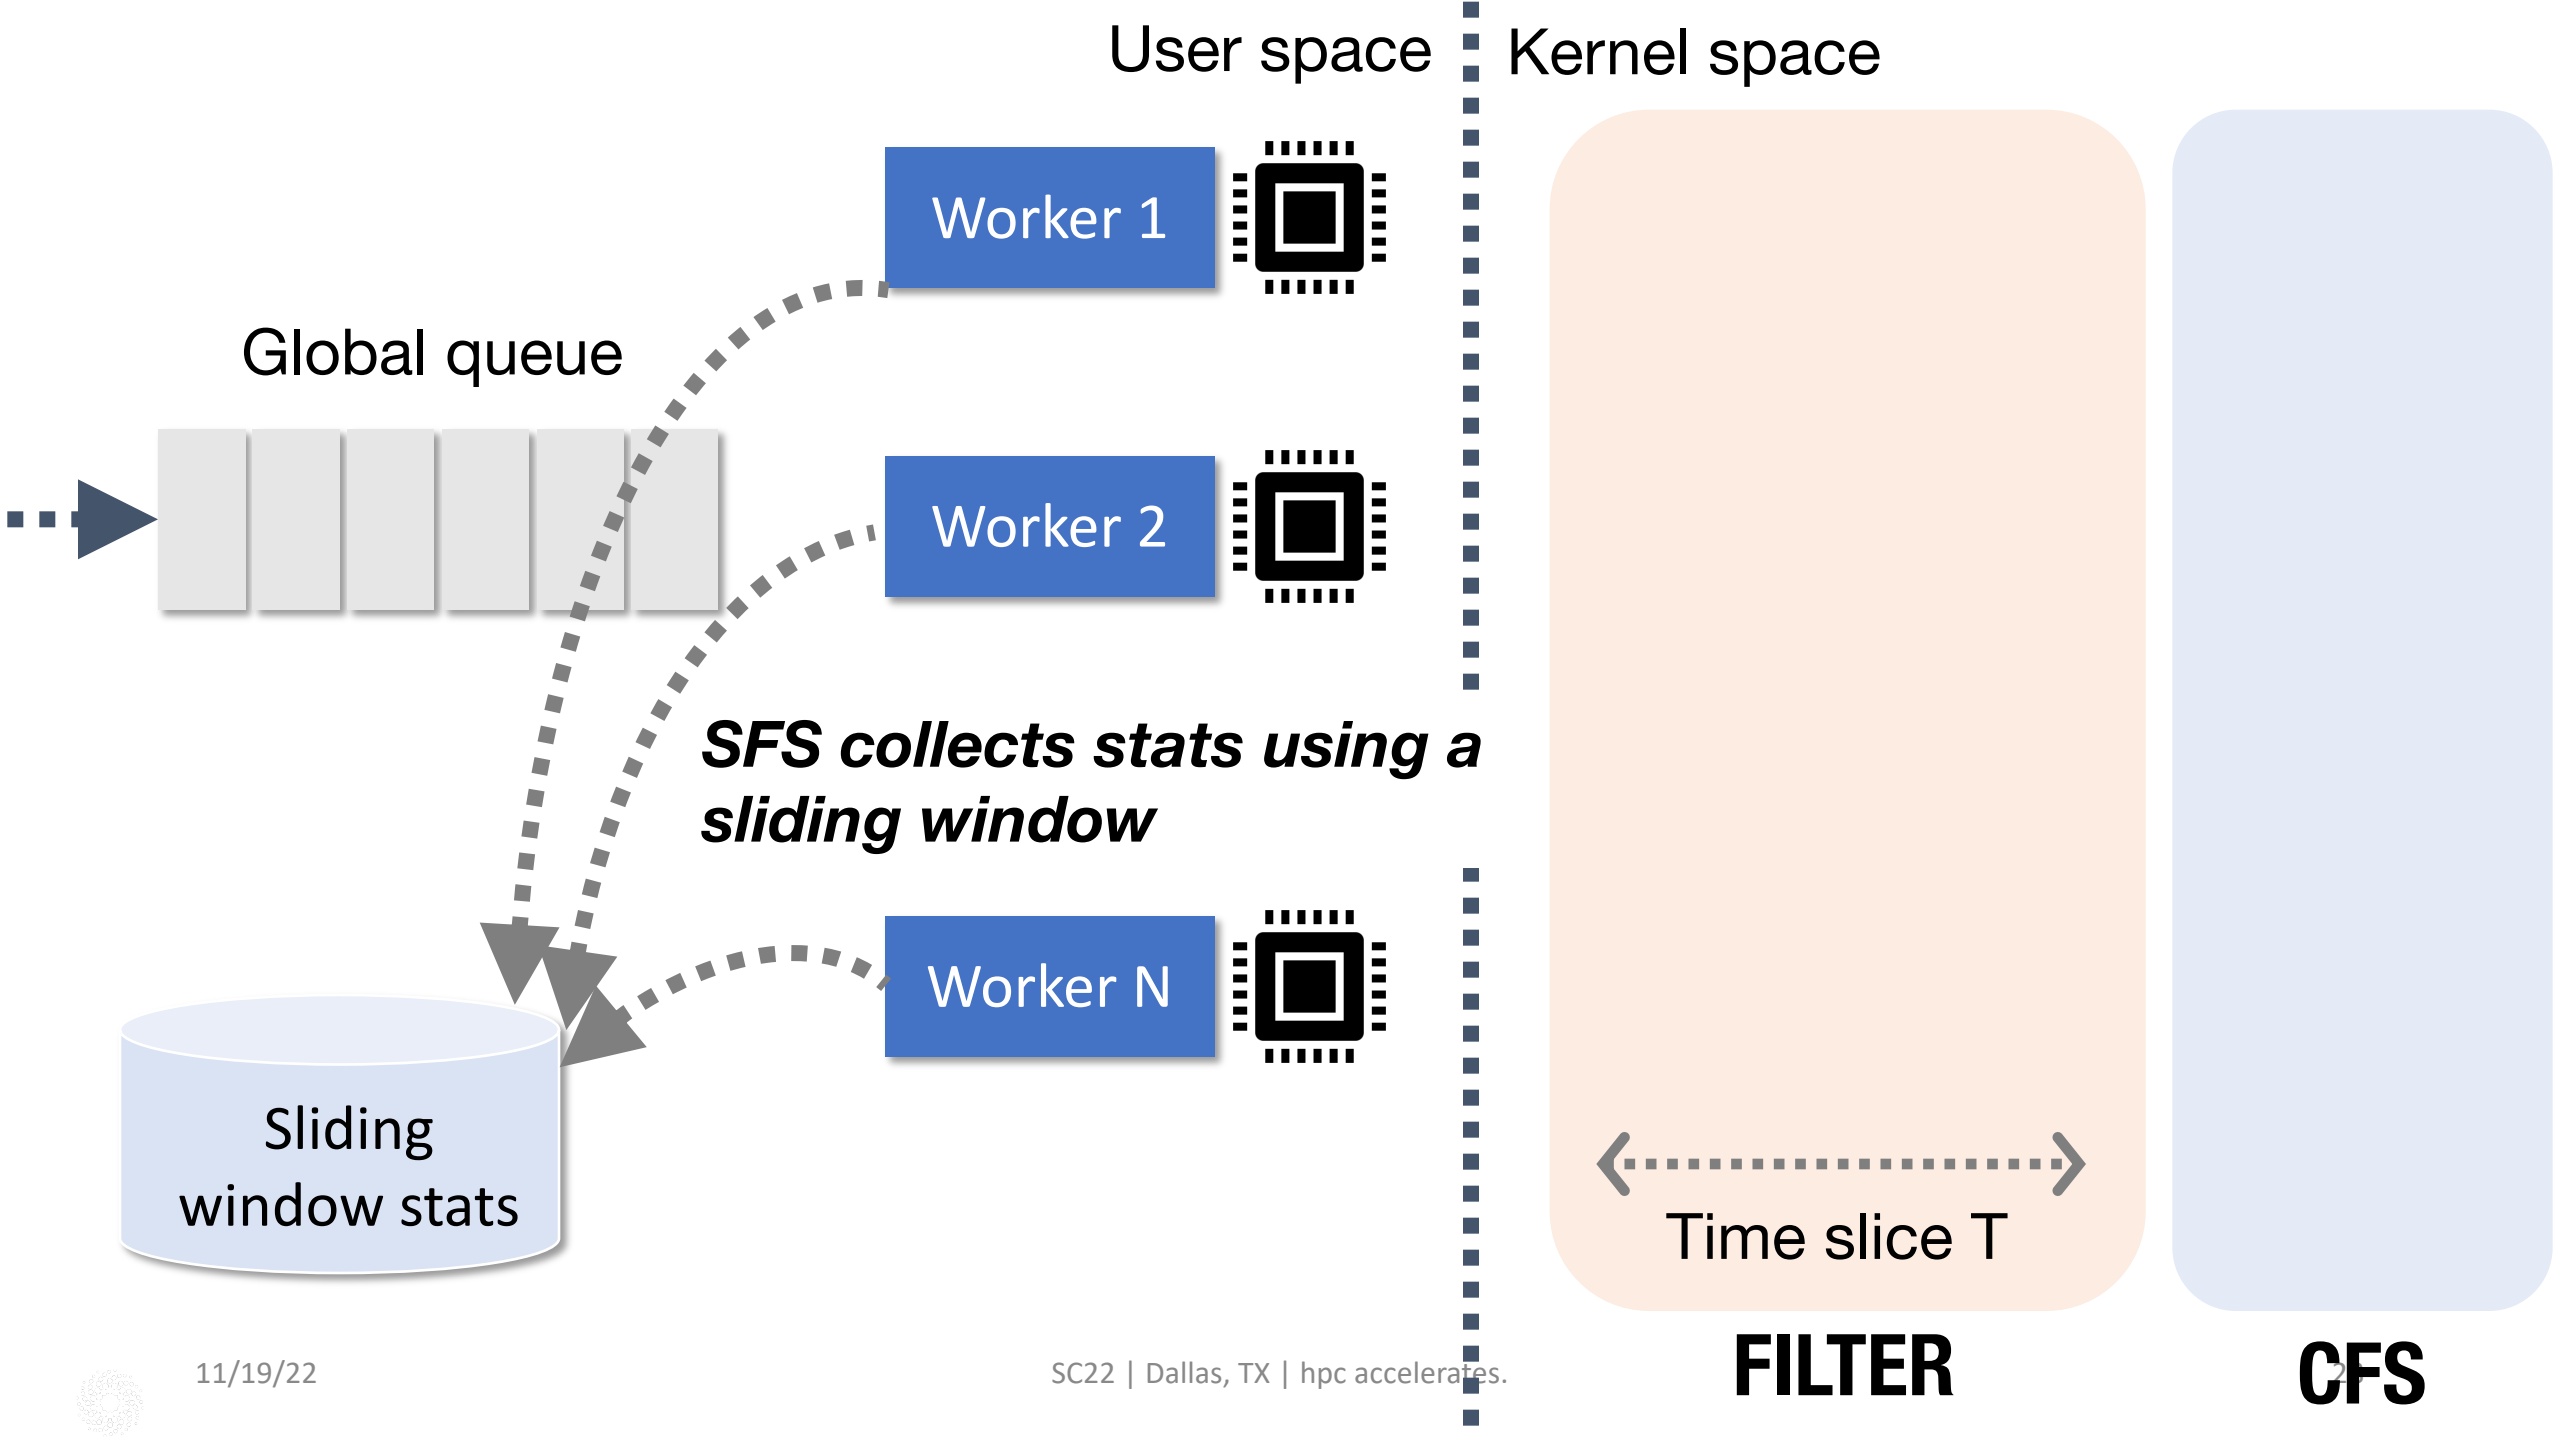


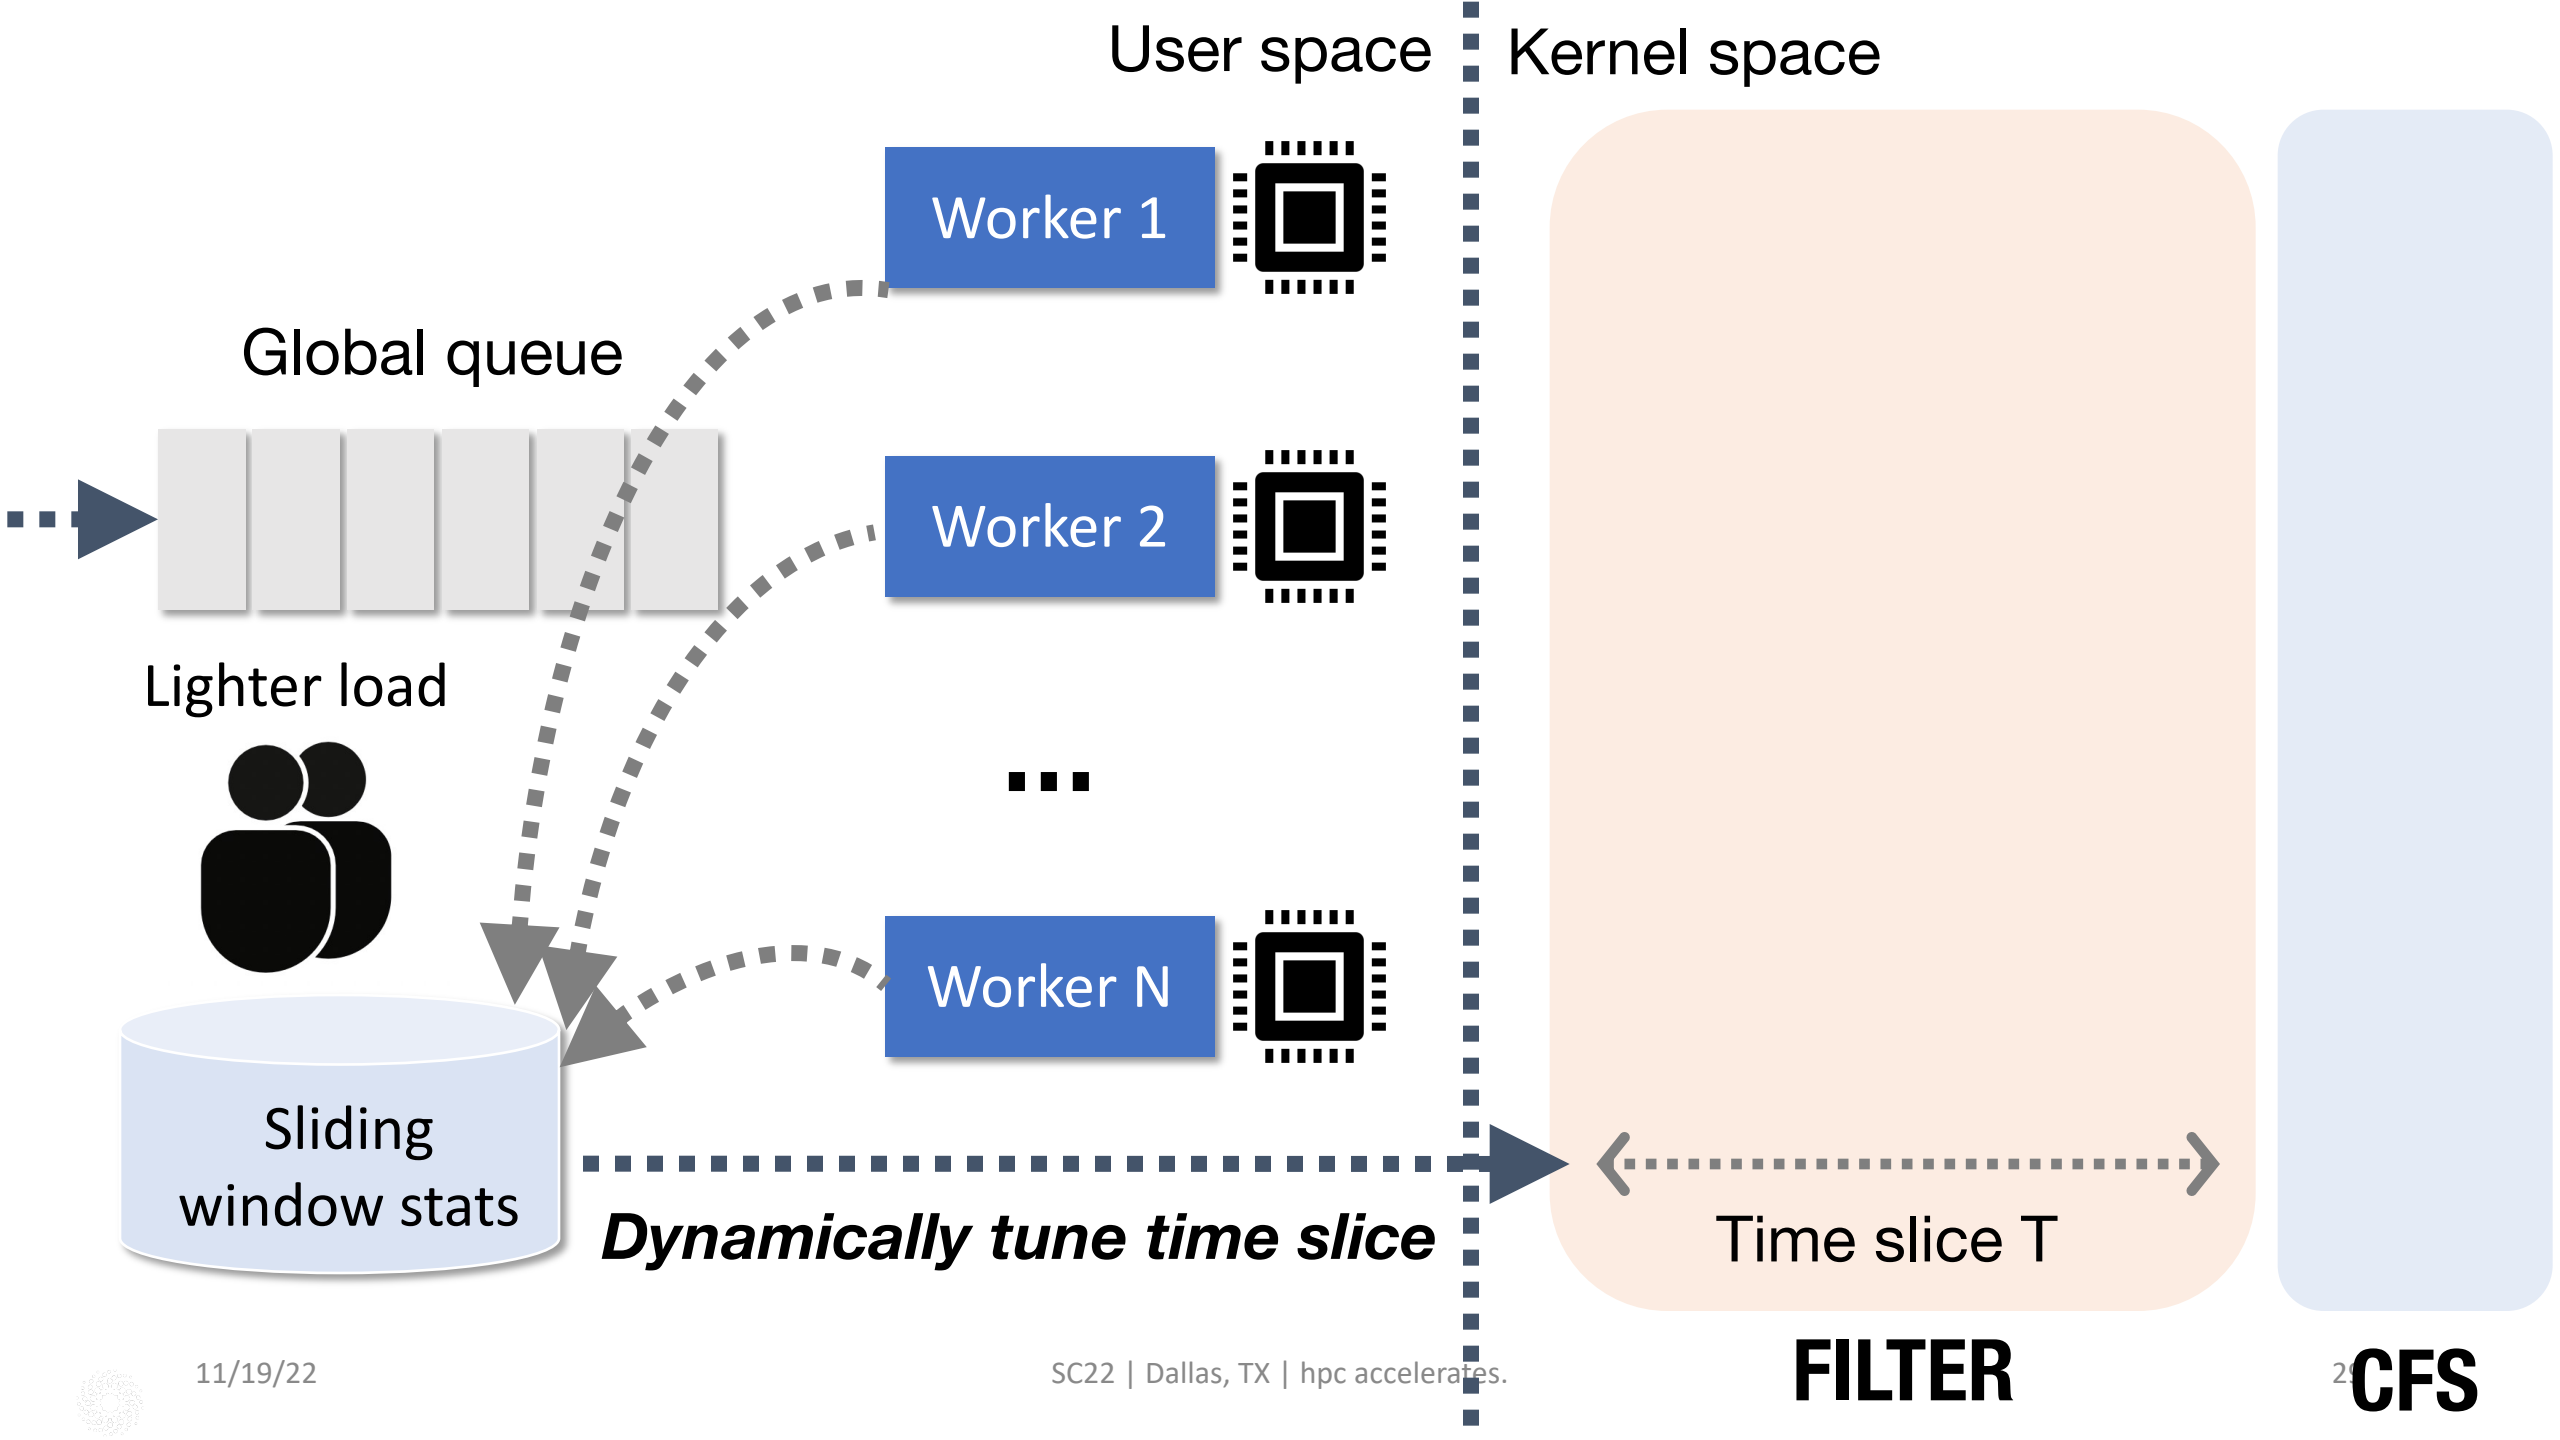


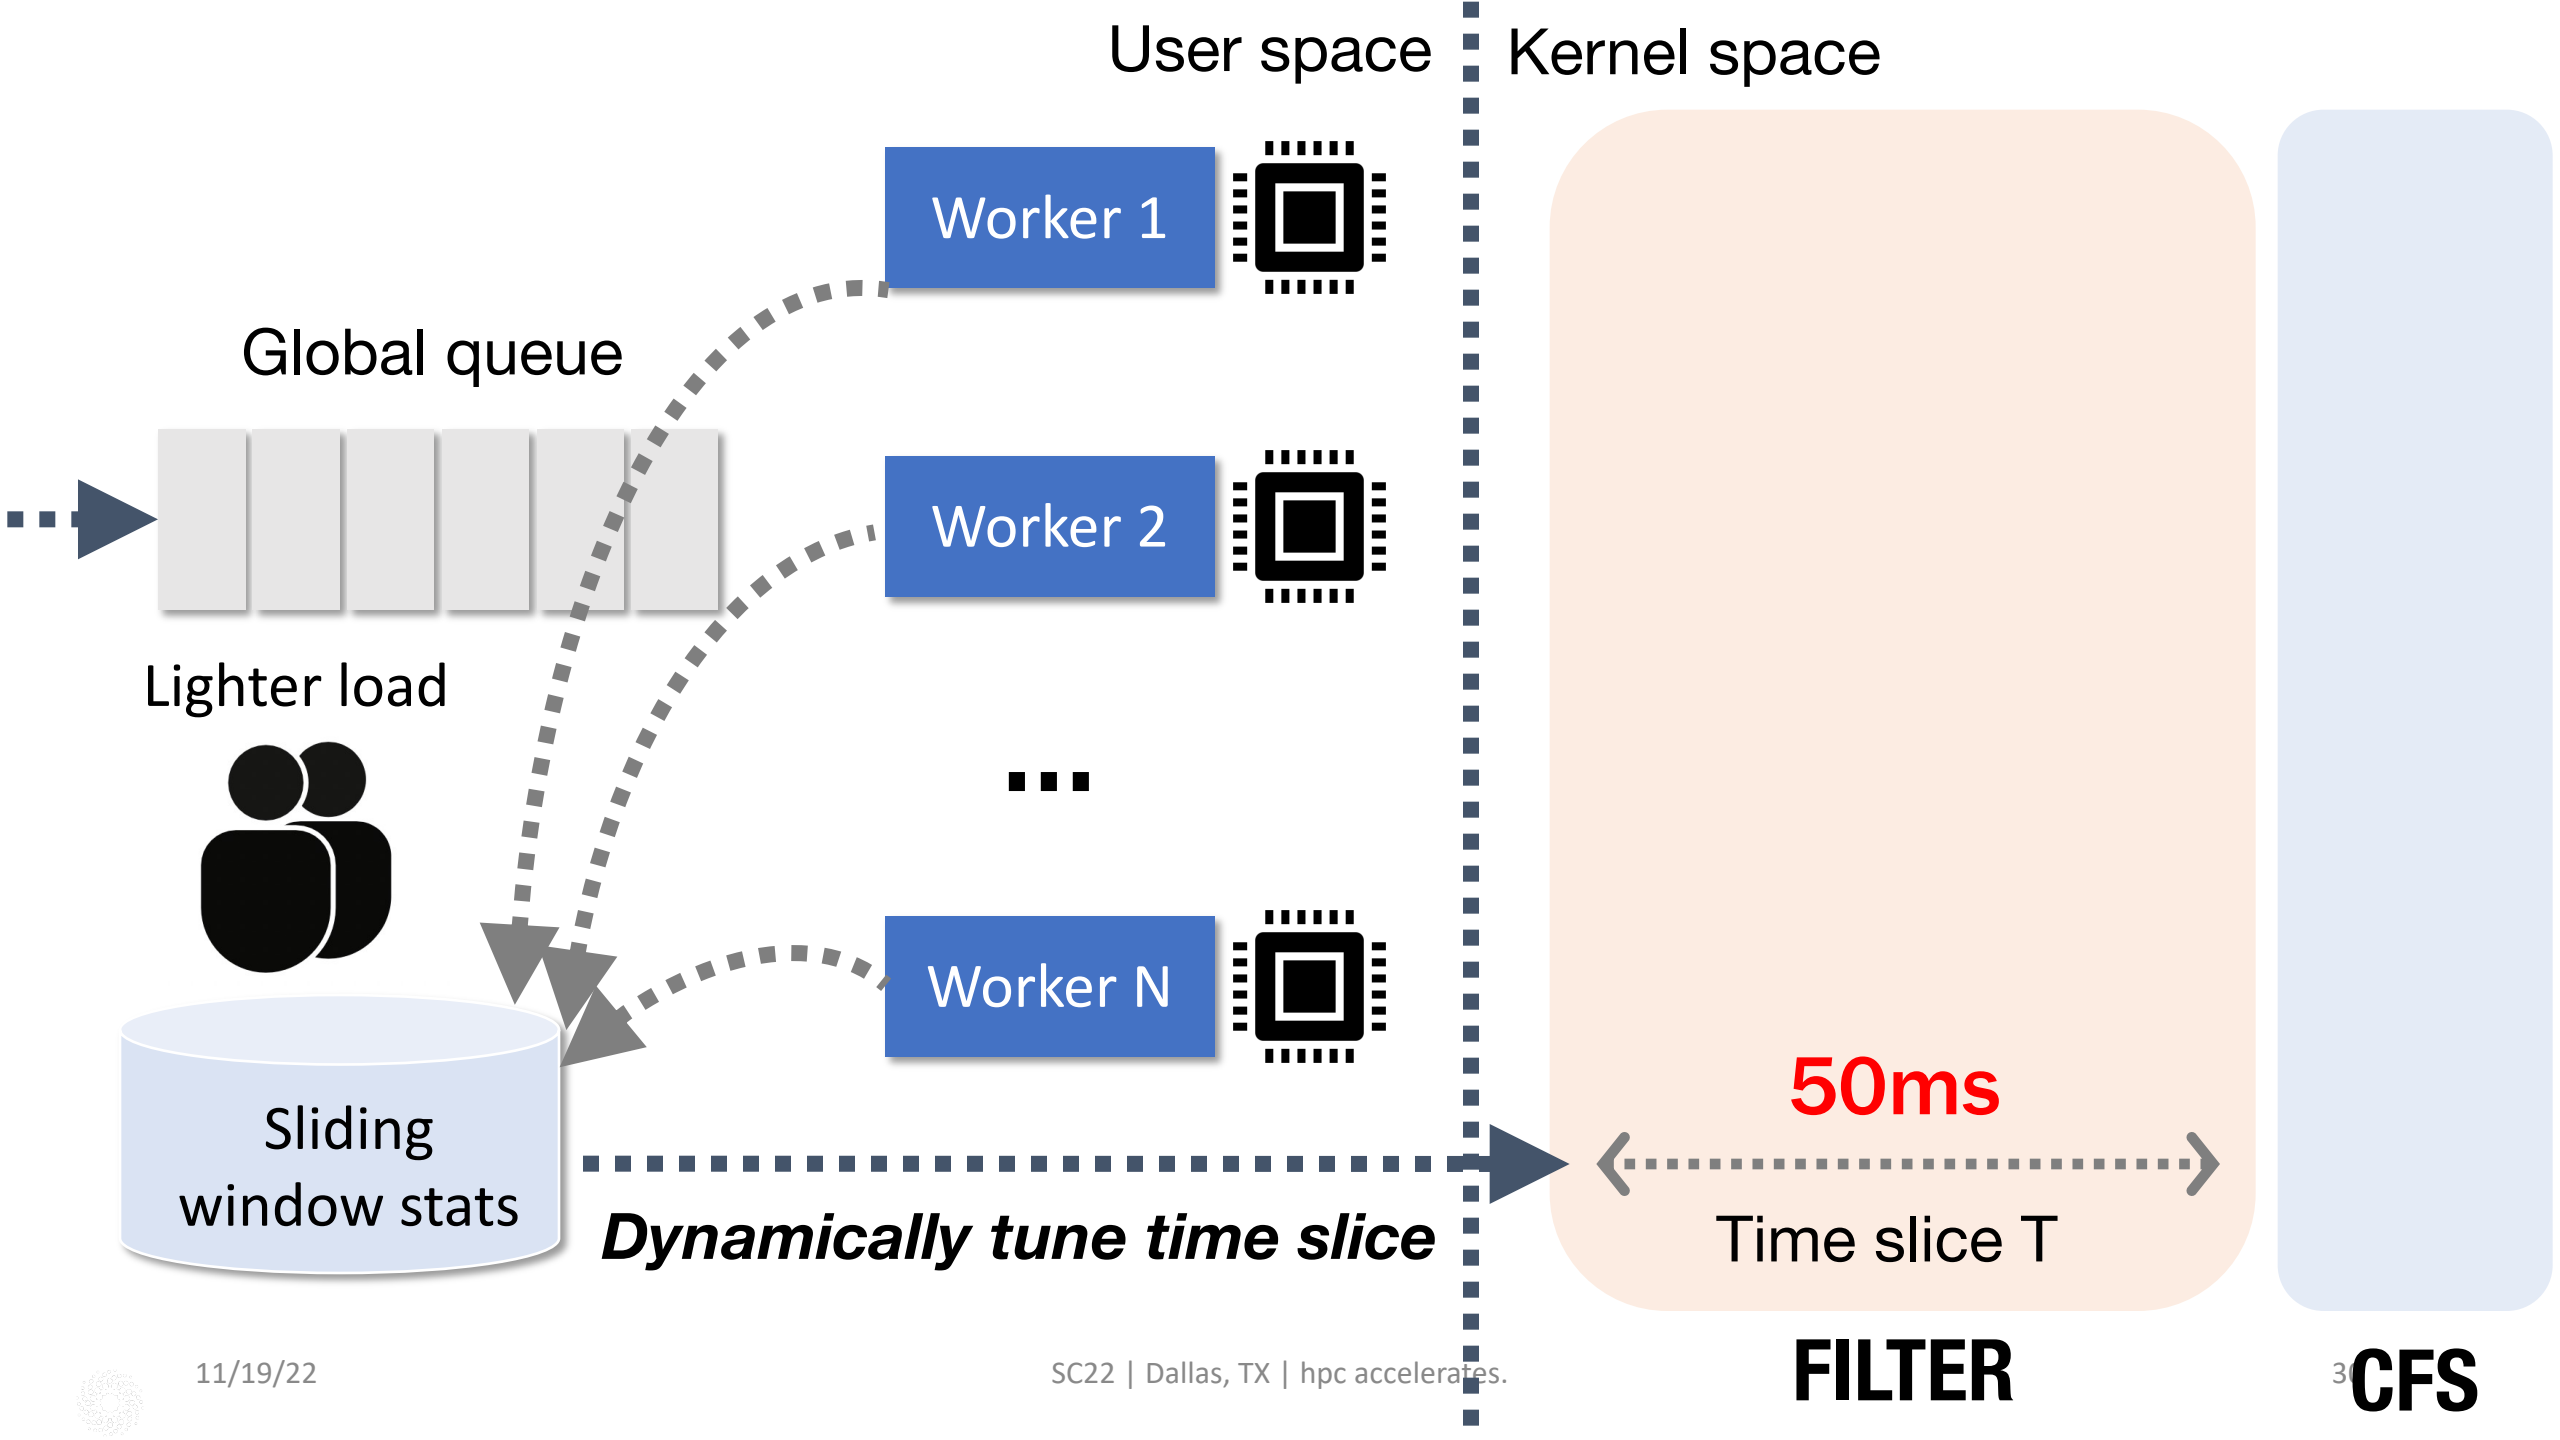


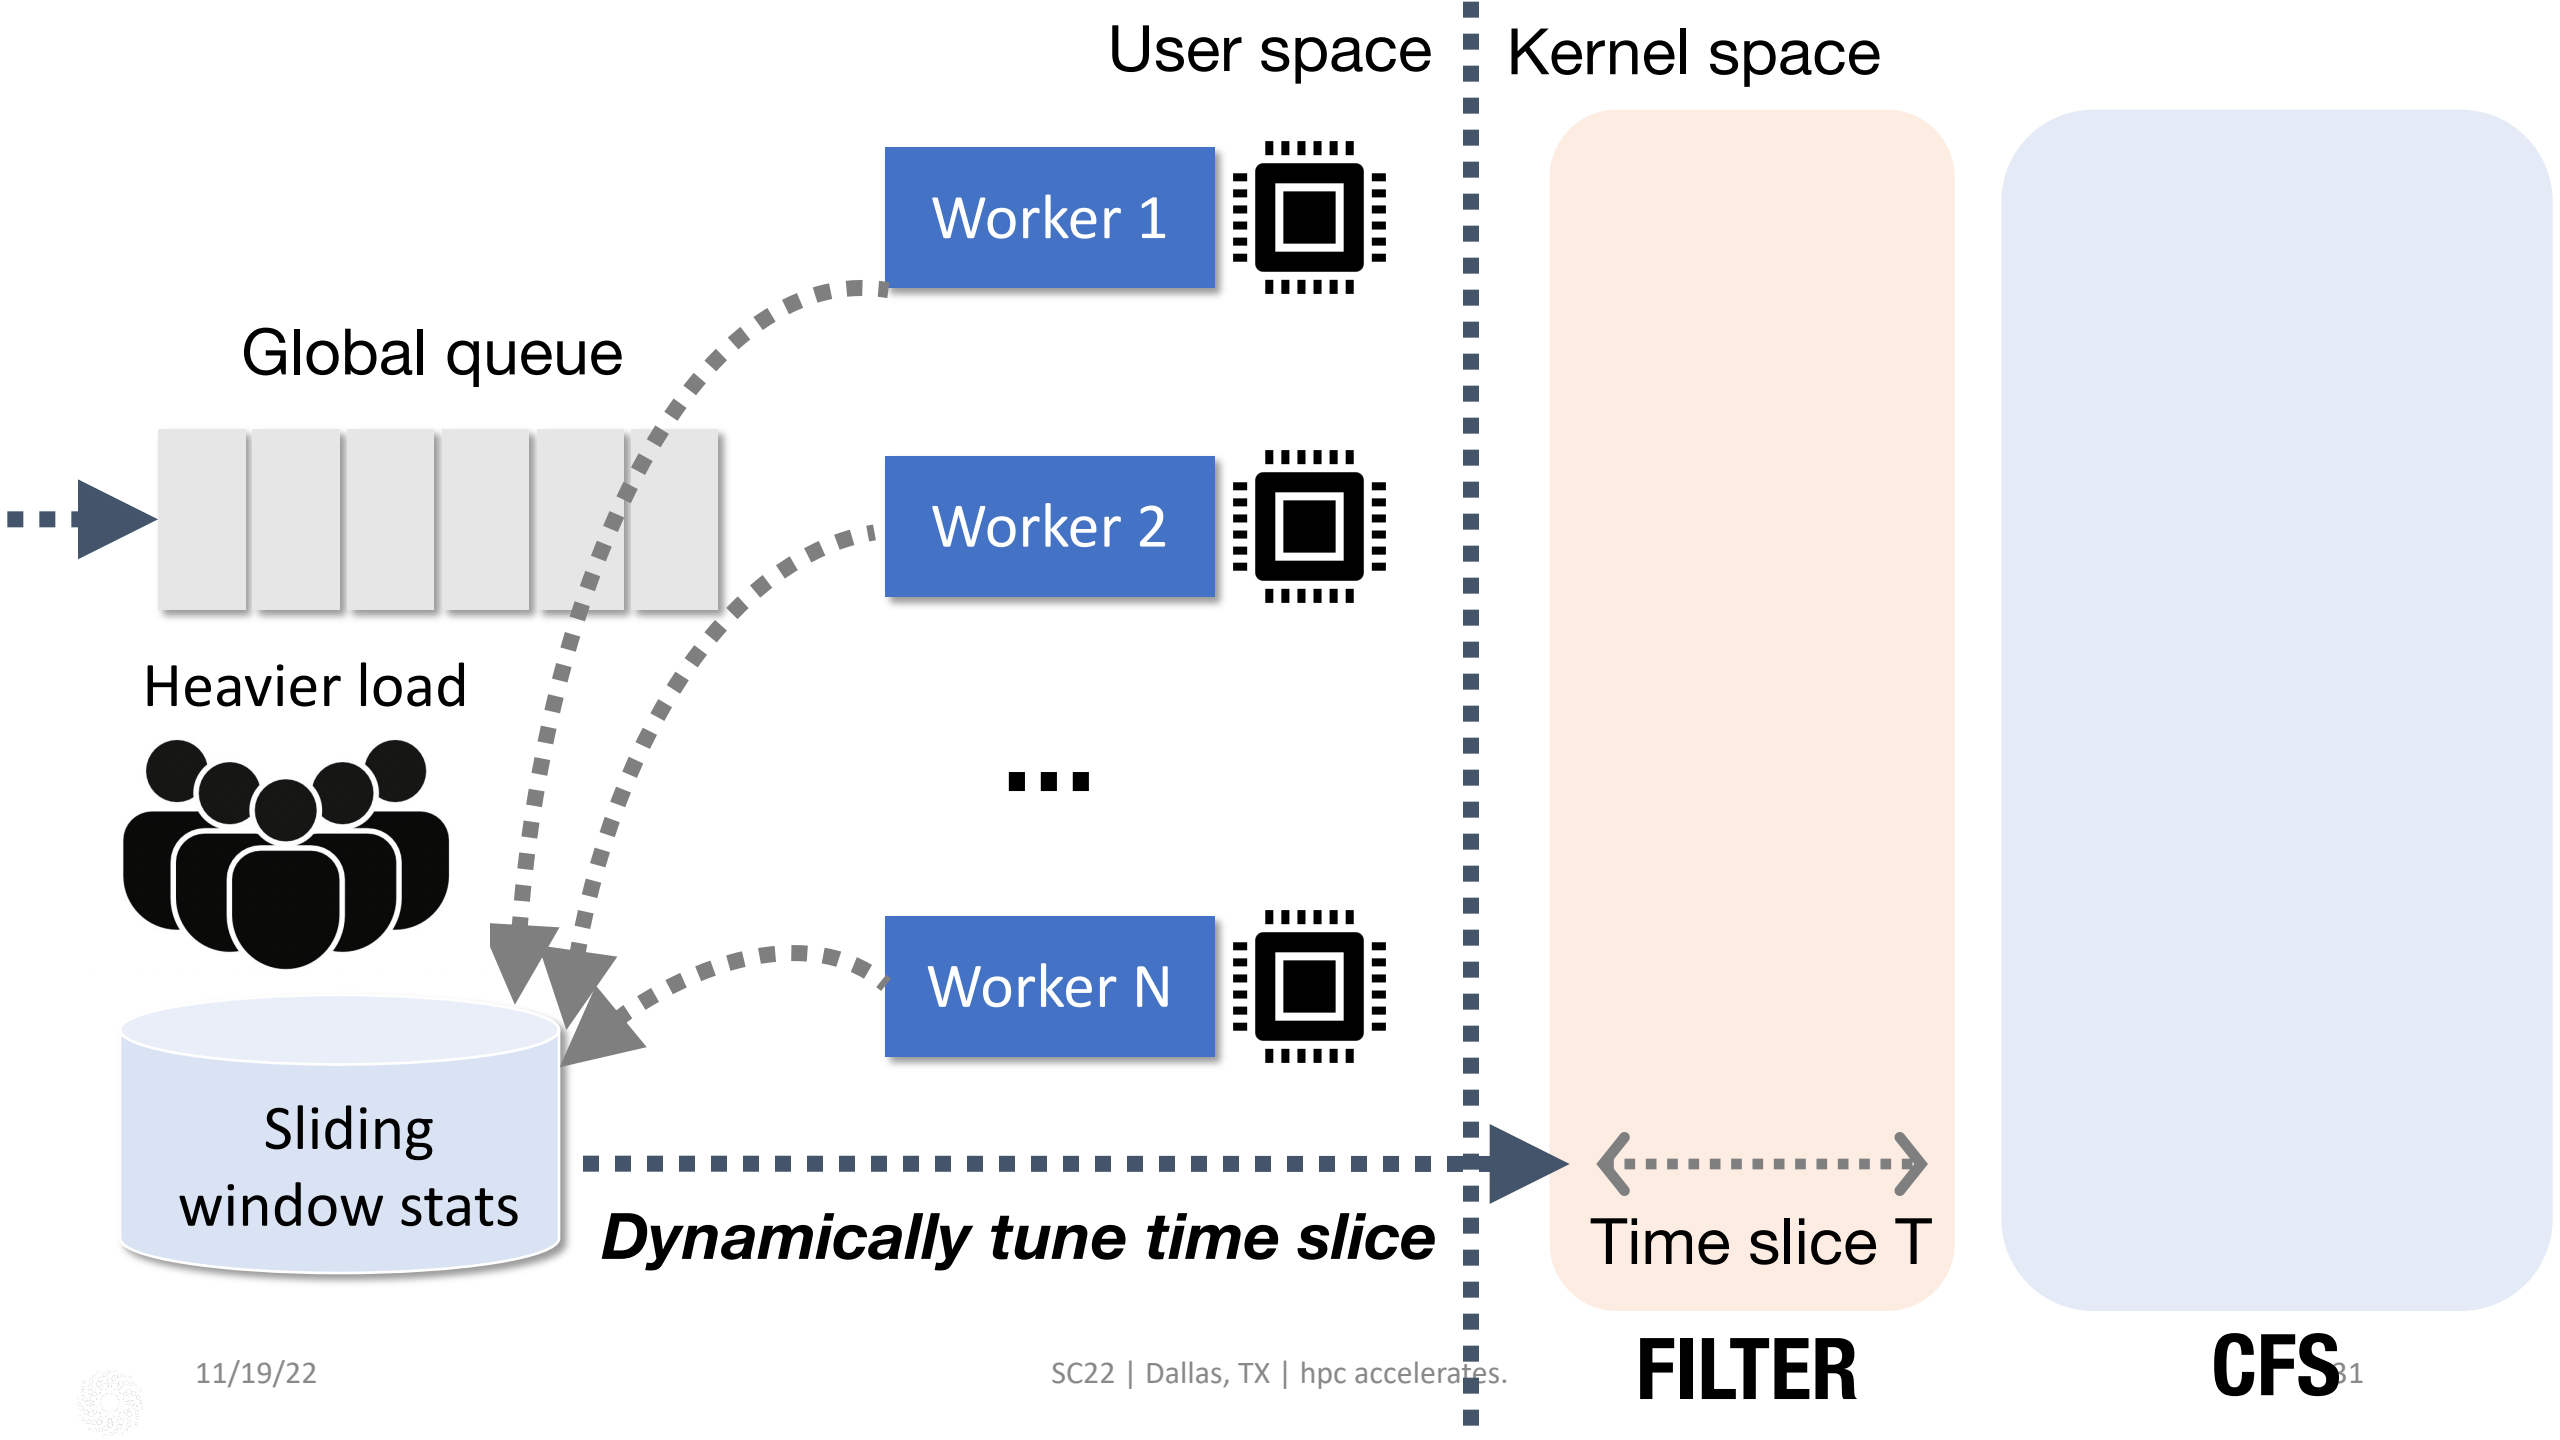


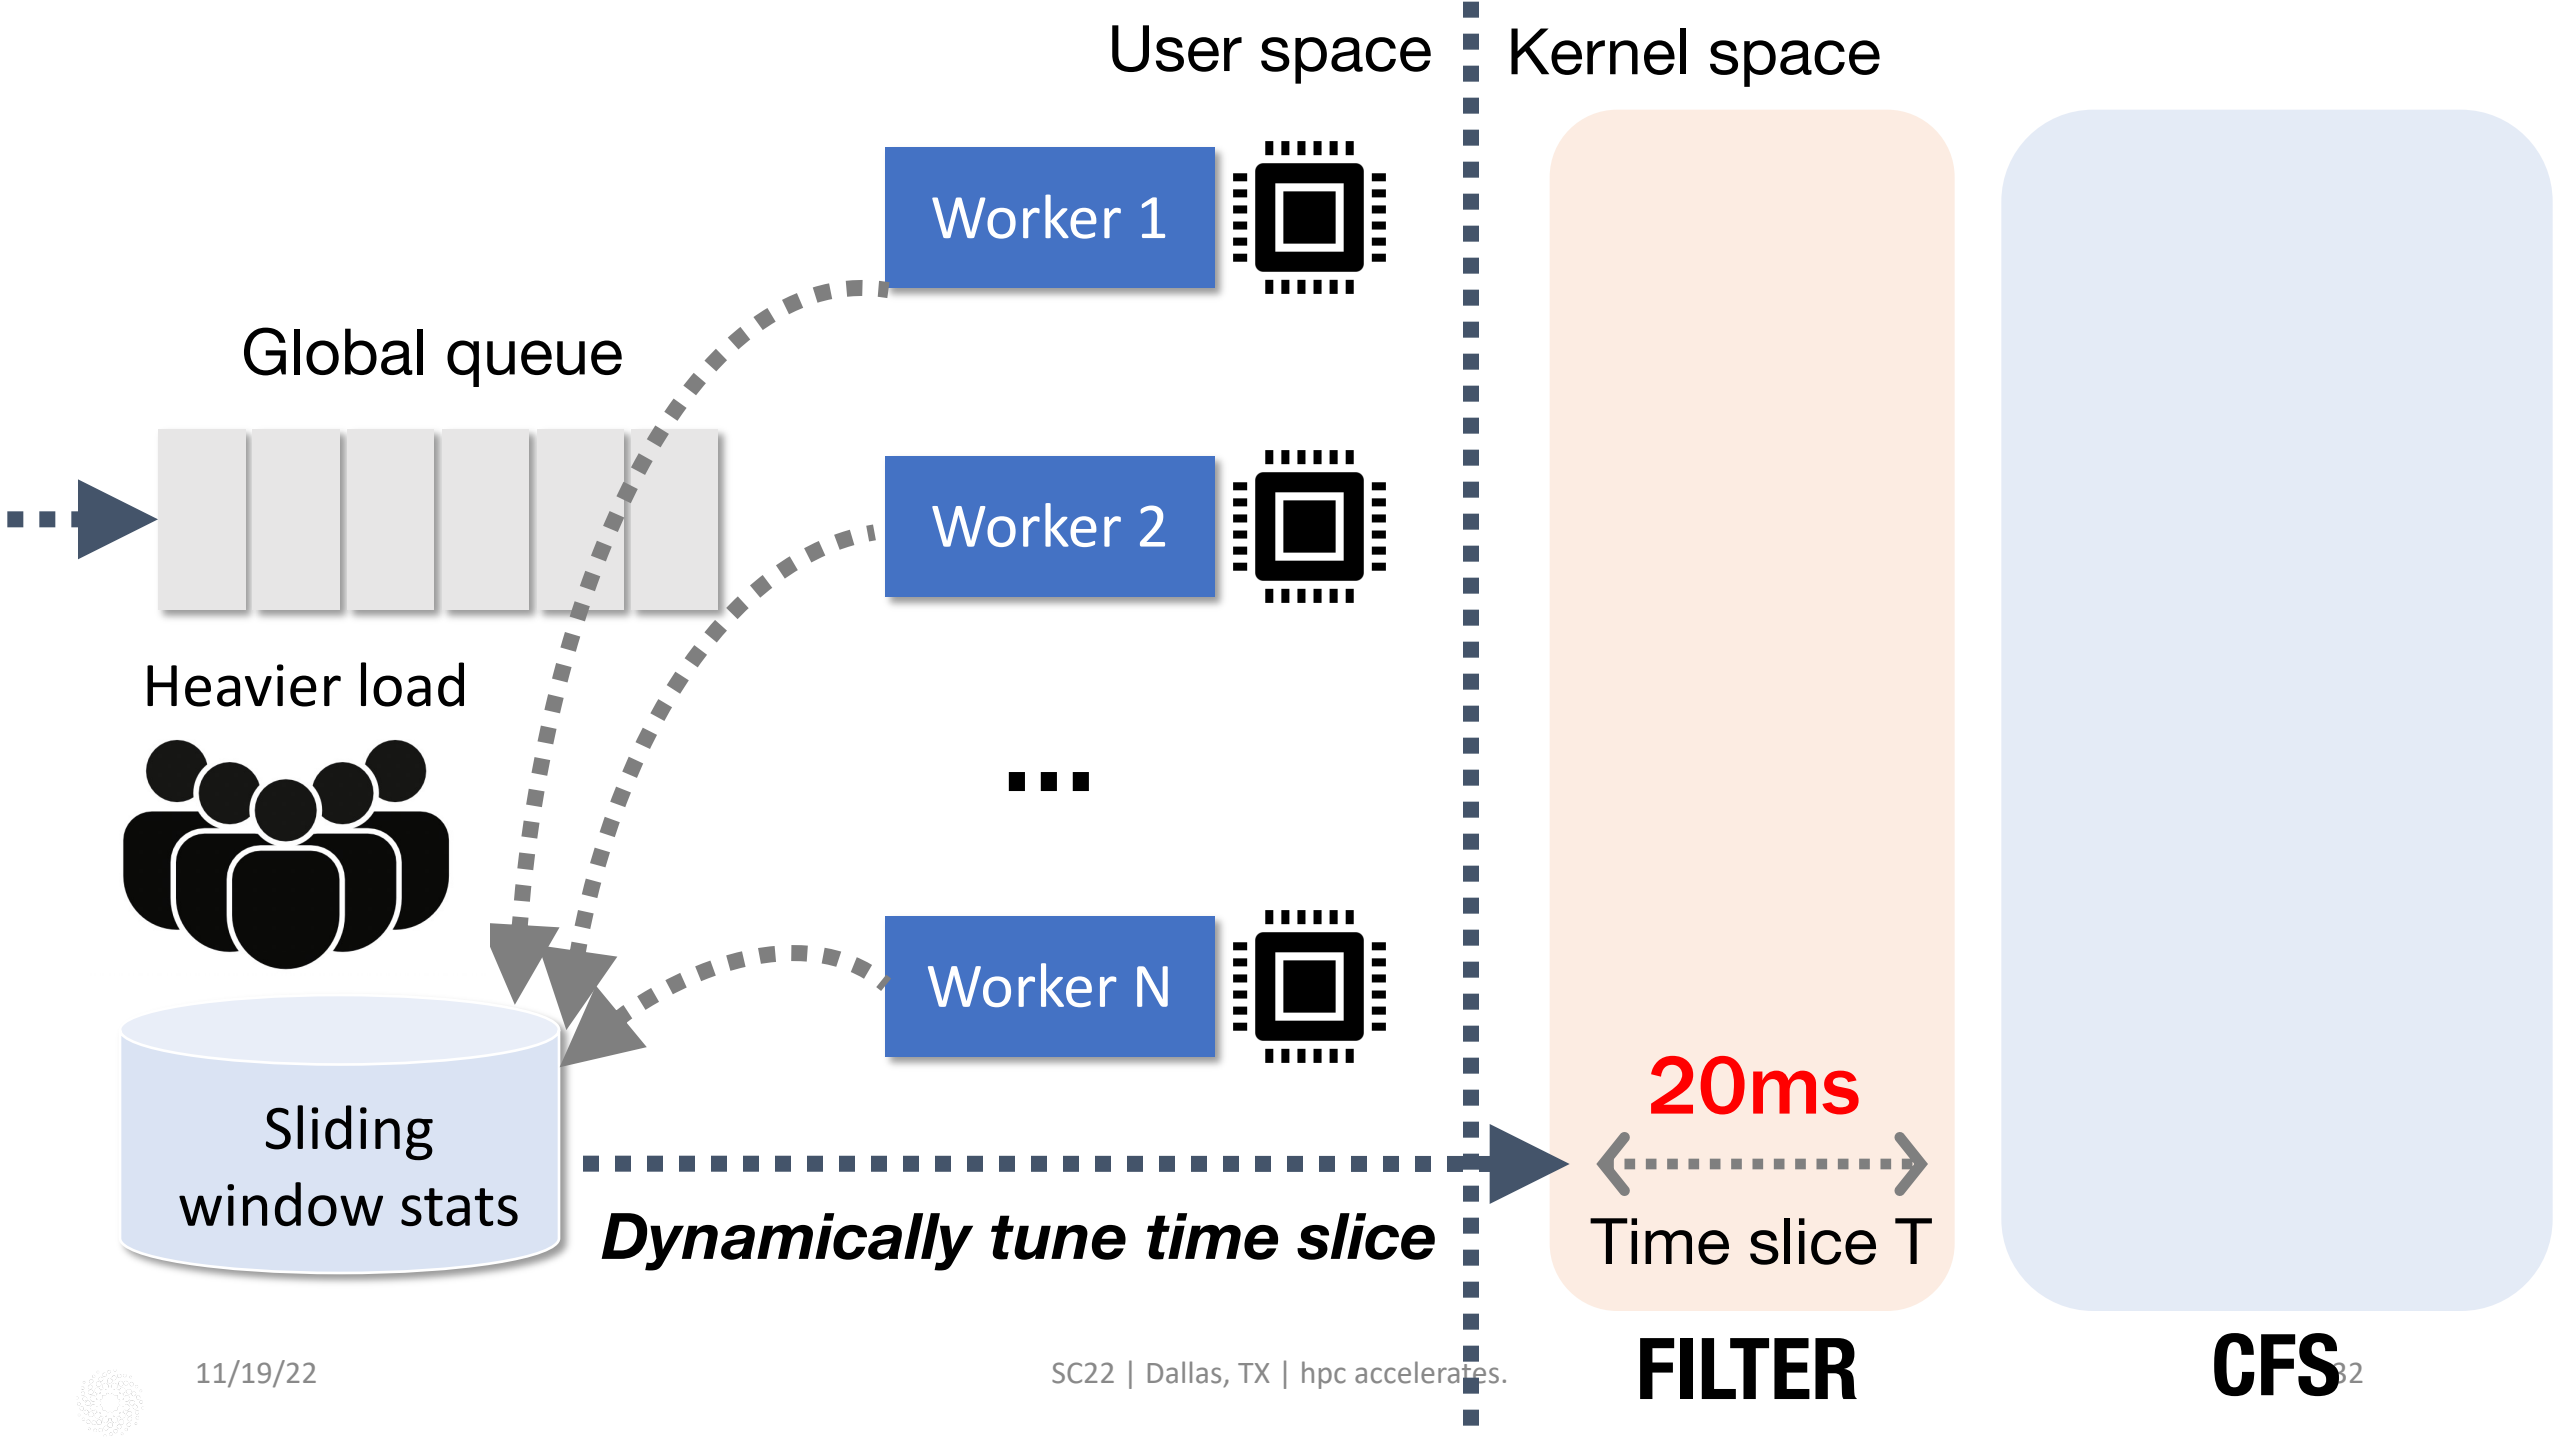




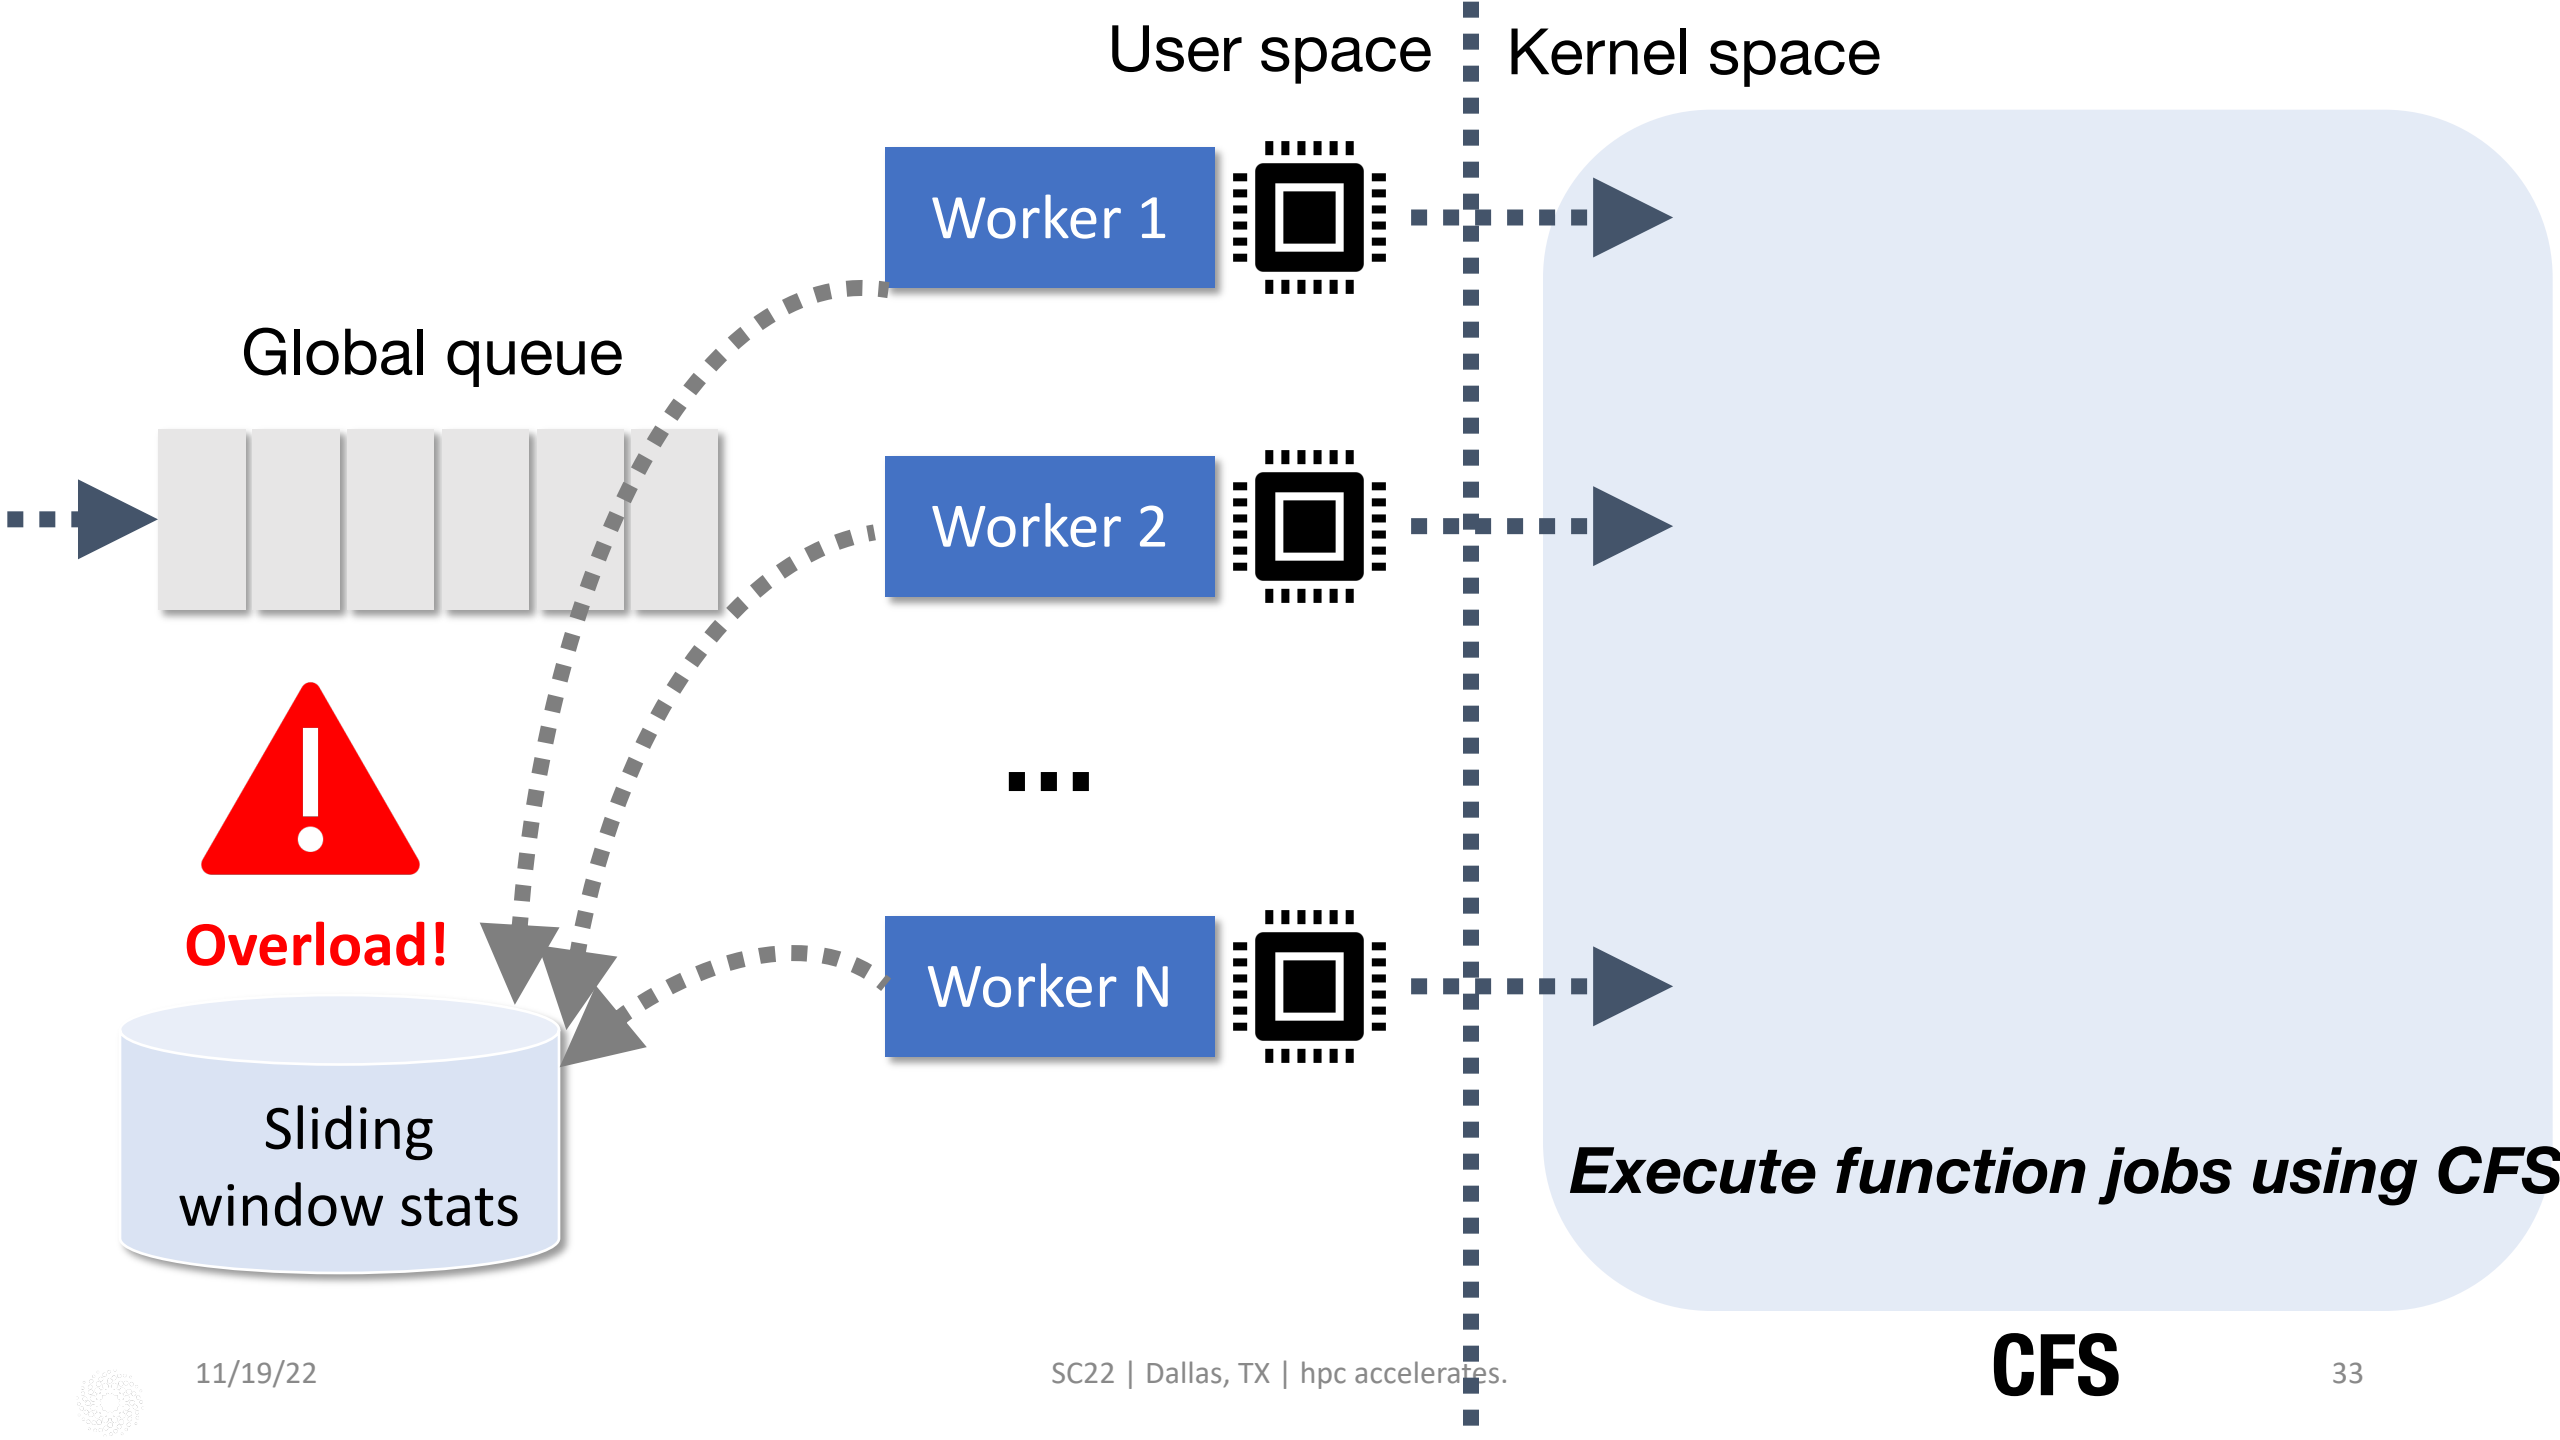












# Smart function scheduler (SFS)

SFS is a **FaaS-aware**, user-space OS scheduler

**Key idea:** Two-level scheduling

- A FILTER level that dynamically tunes a global time slice for newly arrived functions
- Filtered functions from top level continue in CFS



Orchestrates existing Linux scheduling policies

- Short functions **run to completion**
- Online policy with **minimum historical stats**
- **Transparent** to both upper-level FaaS platform and underlying OSES



Improved turnaround time



Simply and practical heuristic



Ready to deploy



# Outline

- SFS Design
- **Evaluation**
- Conclusion




# Experimental Setup

- Standalone
  - **16-core** EC2 VM
- SFS-ported OpenLambda\* (*HotCloud' 16*)
  - **72-core** EC2 bare-metal VM
    - By modifying **29 lines** of Go/Python code in OpenLambda
- Day one of the Azure Functions Trace
  - 49, 712 function requests
  - Breakdowns (min, median, max, percentiles)


\* "Serverless Computation with {OpenLambda}." Hendrickson et al. In *8th USENIX Workshop on Hot Topics in Cloud Computing (HotCloud 16)*. 2016.



# SFS Standalone – Turnaround time




# SFS Standalone – Turnaround time



**SFS maintains almost identical performance for 83% of the function requests across all load levels**

# SFS Standalone – Turnaround time




**Remaining longer  
jobs observed  
slightly higher tail  
latency**

**SFS maintains almost identical performance for 83% of the  
function requests across all load levels**

# SFS Standalone – RTE

Function **Run-Time Effectiveness**  
= Service time / Turnaround time




**SFS performs optimal RTE for short functions**






# SFS-ported OpenLambda



# SFS-ported OpenLambda



**Under 100% load, for more than 85% of functions, CFS suffers 10X more context switches than SFS**

# Conclusion

- SFS addresses the poor performance issue of CFS in FaaS workloads through a **two-level scheduling approach**
- SFS adaptively tunes a high-priority **FILTER** pool that optimizes the performance of short functions
- Experimental results show SFS outperforms CFS up to **50x** for a production FaaS workload



# Thank You Questions?

- Contact: Yuqi Fu (jwx3px@Virginia.edu)
- <https://github.com/ds2-lab/SFS>



National  
Science  
Foundation



CloudLab



# Back slides

